Step 1: Loading in the Data and Initial Plot

```
# --- Setup Chunk ---
knitr::opts_chunk$set(echo = TRUE)
library(tidyverse)
```

```
## Warning: package 'lubridate' was built under R version 4.3.3
```

```
## -- Attaching core tidyverse packages ------------------------ tidyverse 2.0.0 --
## v dplyr     1.1.4     v readr     2.1.5
## v forcats   1.0.0     v stringr   1.5.1
## v ggplot2   3.5.1     v tibble    3.2.1
## v lubridate 1.9.4     v tidyr     1.3.1
## v purrr     1.0.2
## -- Conflicts ------------------------------------------ tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(janitor)
```

```
## Warning: package 'janitor' was built under R version 4.3.3
```

```
##
## Attaching package: 'janitor'
##
## The following objects are masked from 'package:stats':
##
##     chisq.test, fisher.test
```

```
library(lubridate)
library(MASS)
```

```
##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##     select
```

```
library(car)
```

```
## Loading required package: carData
##
## Attaching package: 'car'
##
## The following object is masked from 'package:dplyr':
##
##     recode
##
## The following object is masked from 'package:purrr':
##
##     some
```

```
library(randtests)
```

```
## Warning: package 'randtests' was built under R version 4.3.3
```

```
library(forecast) # Needed for ARIMA/HW
```

```
## Warning: package 'forecast' was built under R version 4.3.3
```

```
## Registered S3 method overwritten by 'quantmod':
##   method              from
##   as.zoo.data.frame zoo
```

```
# --- Load Data ---
bike <- read.csv("trips_per_day.csv")
bike <- bike %>% filter(!is.na(trip_date))
bike$trip_date <- as.Date(bike$trip_date)

# Filter for 2017-2024 (Project Scope)
bike <- bike[bike$trip_date >= as.Date("2017-01-01"), ]

# Create Time Series Object (Daily Frequency)
y <- bike$n_trips
bike_ts <- ts(y, start = c(2017, 1), frequency = 365)

# Initial Plot
plot(bike_ts, main = "Daily BikeShare Trips (2017-2024)", ylab = "Trips")
```
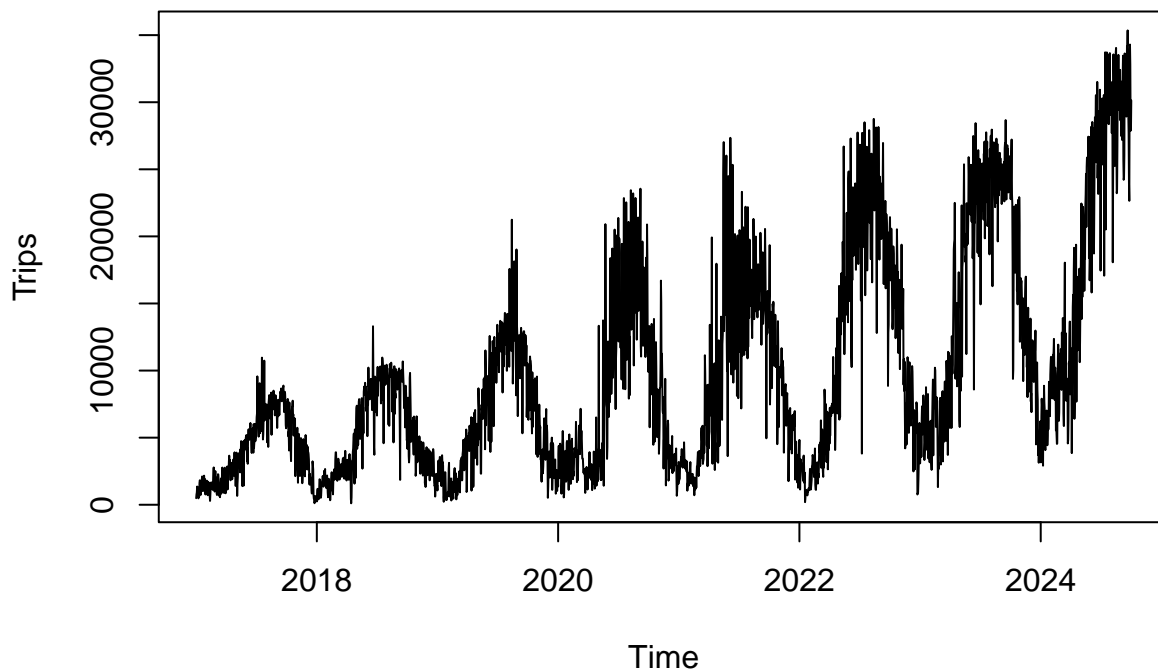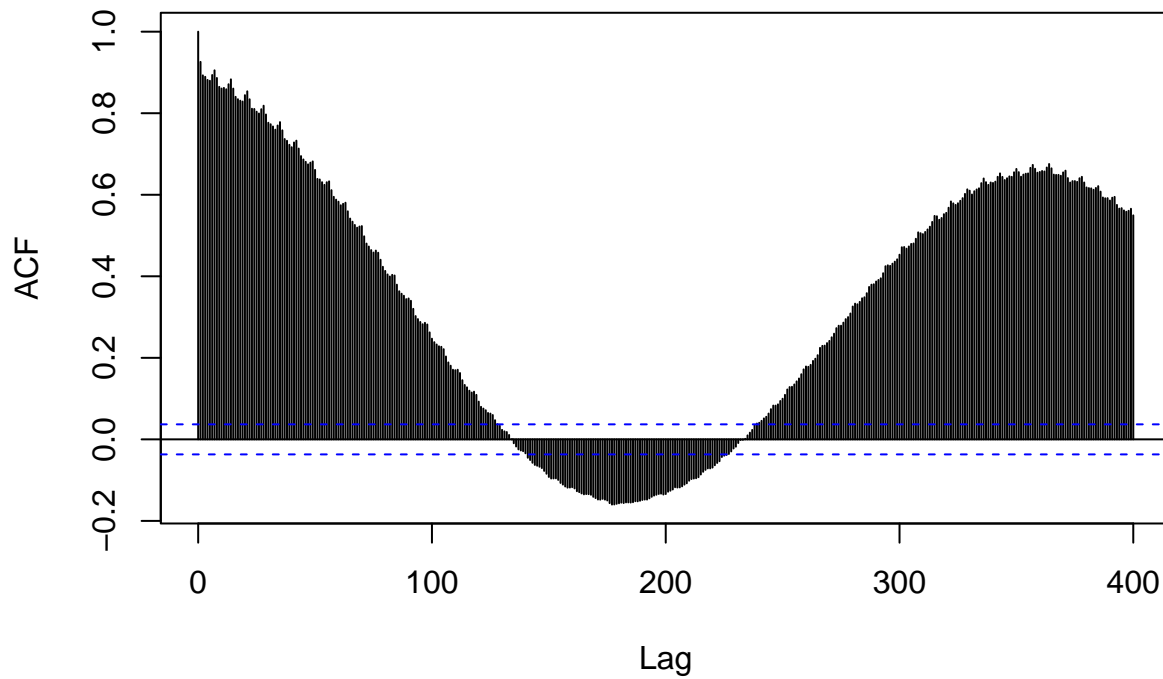
Daily BikeShare Trips (2017–2024)

Step 2: Checking ACF and Period

```
acf(y, lag.max = 400, main = "ACF of Raw Data")
```

2

ACF of Raw Data


```r
spec <- spectrum(y, spans = 5, main = "Spectral Density")
```
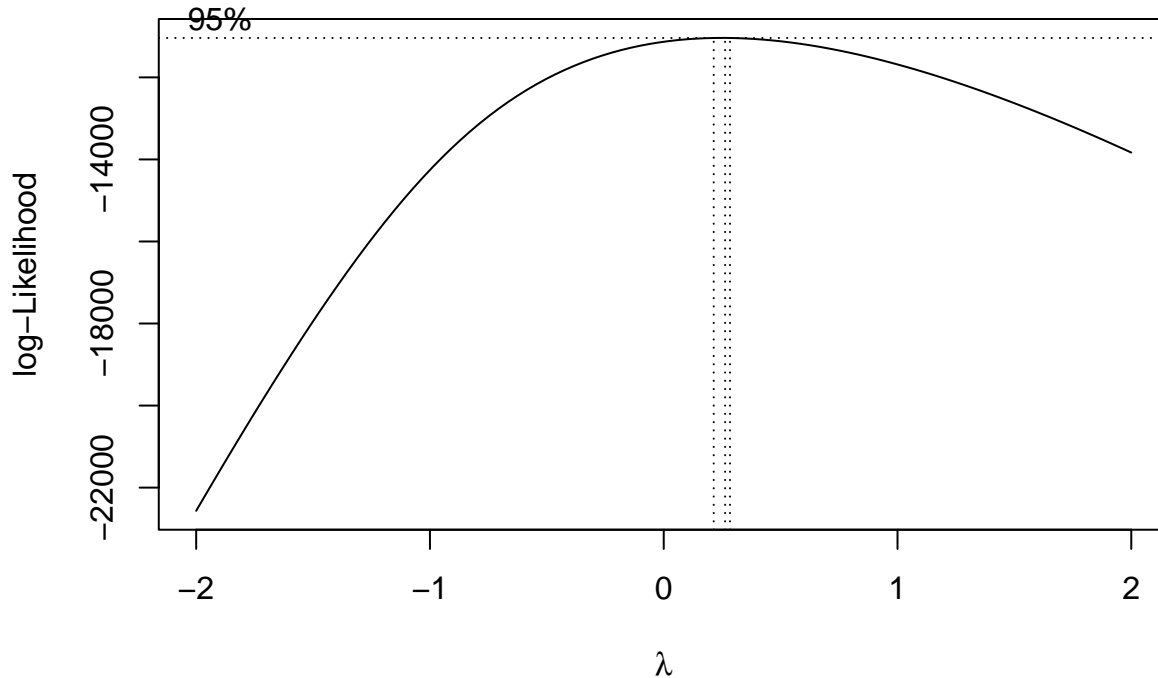
Spectral Density


```r
# Check dominant period
print(1 / spec$freq[which.max(spec$spec)])
```

```
## [1] 360
```

Step 3: Find Lambda and Do BoxCox Transformation

```
# Box-Cox Check
bc <- boxcox(lm(y ~ 1), lambda = seq(-2, 2, 0.1))
```



```
lam <- bc$x[which.max(bc$y)]
print(paste("Optimal Lambda:", lam))
```

```
## [1] "Optimal Lambda: 0.262626262626263"
```

```
if (lam == 0) {
  y_trans <- log(bike_ts)
} else if (lam > 0) {
  y_trans <- (bike_ts^lam - 1) / lam
} else {
  # negative lambda → use minus sign trick like in lectures
  y_trans <- -(bike_ts^lam)
}

plot(y_trans, main = "Transformed Daily Trips", ylab = "Transformed trips")
```

Transformed Daily Trips


```r
bike$y_trans <- y_trans # Add to dataframe for regression
```

Step 4: STL Decomposition

```r
decomp <- stl(ts(y_trans, frequency=365), s.window="periodic")
plot(decomp, main="Decomposition of Transformed Data")
```

Decomposition of Transformed Data

Step 5: Mutating the Bike Object to add Seasonality Factors

```r
bike <- bike %>%
  mutate(
    # 1. Time Index
    time_index = 1:n(),

    # 2. Factor-based Seasonality
    month_fac = as.factor(month(trip_date)),
    weekday_fac = factor(wday(trip_date, label = TRUE), ordered = FALSE),
    is_weekend = as.factor(ifelse(wday(trip_date) %in% c(1, 7), "Yes", "No")),

    # 3. Fourier Terms for Smooth Annual Seasonality
    # Using period = 365.25 to account for leap years over the long dataset
    sin_year = sin((2 * pi * time_index / 360) + 0.75),
  )

# Train on 2017-2022, Validate on 2023
train_seas <- subset(bike, trip_date < "2023-01-01")
valid_seas <- subset(bike, trip_date >= "2023-01-01" & trip_date < "2024-01-01")
```

Step 6: Trying out different models

```r
# --- Model Fitting ---

# 1. Trend + Fourier Only
mod1 <- lm(y_trans ~ time_index + sin_year , data = train_seas)

# 2. Trend + Fourier + Month + Day of Week
mod2 <- lm(y_trans ~ time_index + sin_year + weekday_fac, data = train_seas)
```

```r
# 3. Trend + Fourier + Month
mod3 <- lm(y_trans ~ time_index + sin_year  + month_fac, data = train_seas)

# 4. Trend + Fourier + Month + Weekend/Weekday
mod4 <- lm(y_trans ~ time_index + sin_year  + month_fac + is_weekend, data = train_seas)

# 5. Trend + Fourier + Month + Day of Week
mod5 <- lm(y_trans ~ time_index + sin_year  + month_fac + weekday_fac, data = train_seas)


# --- Model Evaluation ---

# Function to calculate evaluation metrics
evaluate_model <- function(model, train_df, valid_df) {
  # Calculate predictions on validation set
  preds <- predict(model, newdata = valid_df)

  # Calculate APSE (Mean Squared Error on Validation Set)
  apse <- mean((valid_df$y_trans - preds)^2)

  # Number of parameters (k): coefficients (beta) + 1 (sigma)
  num_params <- length(coef(model)) + 1

  # Extract AICc (using manual calculation for standard lm objects)
  n <- nrow(train_df)
  aicc <- AIC(model) + (2 * num_params * (num_params + 1)) / (n - num_params - 1)

  # Extract Adjusted R-squared
  adj_r2 <- summary(model)$adj.r.squared

  # Return metrics including Number of Parameters
  # We return just the number of coefficients for clarity as "Model Size"
  n_coeffs <- length(coef(model))

  return(c(APSE = apse, AICc = aicc, Adj_R2 = adj_r2, Num_Params = n_coeffs))
}

# Collect results
results <- rbind(
  "Trend + Fourier" = evaluate_model(mod1, train_seas, valid_seas),
  "Trend + Fourier + DayOfWeek" = evaluate_model(mod2, train_seas, valid_seas),
  "Trend + Fourier + Month" = evaluate_model(mod3, train_seas, valid_seas),
  "Trend + Fourier + Month + Weekend" = evaluate_model(mod4, train_seas, valid_seas),
  "Trend + Fourier + Month + DayOfWeek" = evaluate_model(mod5, train_seas, valid_seas)
)

print("Model Evaluation Results:")
```

```
## [1] "Model Evaluation Results:"
```

```r
print(results)
```

```
##                                      APSE      AICc    Adj_R2 Num_Params
## Trend + Fourier                   17.04371 12342.44 0.7780830          3
```

```
## Trend + Fourier + DayOfWeek               16.55910 12299.21 0.7830208            9
## Trend + Fourier + Month                   14.36986 12246.10 0.7887089           14
## Trend + Fourier + Month + Weekend         14.09269 12226.84 0.7906552           15
## Trend + Fourier + Month + DayOfWeek       13.90702 12200.48 0.7936434           20
```

```r
# --- Visualizing Fits on Full Data (2017-Present) ---

# Create a dataframe with all data for plotting
# We use the 'bike' dataframe which contains everything (Train + Valid + Test)
plot_data <- bike %>%
  dplyr::select(trip_date, y_trans, time_index, month_fac, weekday_fac, is_weekend, sin_year)

# Generate predictions for the entire timeline for each model
# Note: We are predicting using the models trained ONLY on 2017-2022 data
# This shows how well the training generalizes to the future.
plot_data$Pred_Mod1 <- predict(mod1, newdata = plot_data)
plot_data$Pred_Mod2 <- predict(mod2, newdata = plot_data)
plot_data$Pred_Mod3 <- predict(mod3, newdata = plot_data)
plot_data$Pred_Mod4 <- predict(mod4, newdata = plot_data)
plot_data$Pred_Mod5 <- predict(mod5, newdata = plot_data)

# Reshape for ggplot
plot_long <- plot_data %>%
  dplyr::select(trip_date, y_trans, starts_with("Pred")) %>%  # <--- ADD dplyr:: HERE
  pivot_longer(cols = c(y_trans, starts_with("Pred")),
               names_to = "Series",
               values_to = "Value")

# Plot
ggplot(plot_long, aes(x = trip_date, y = Value, color = Series)) +
  geom_line(alpha = 0.6) +
  # Highlight the training cutoff
  geom_vline(xintercept = as.Date("2023-01-01"), linetype = "dashed", color = "black") +
  labs(title = "Model Fits vs. Actual Data (2017-2024)",
       subtitle = "Models trained on data before 2023 (left of dashed line)",
       y = "Transformed Trips", x = "Date") +
  theme_minimal() +
  scale_color_manual(values = c("y_trans" = "gray",
                                "Pred_Mod1" = "orange",
                                "Pred_Mod2" = "blue",
                                "Pred_Mod3" = "green",
                                "Pred_Mod4" = "red",
                                "Pred_Mod5" = "purple")) +
  theme(legend.position = "bottom")
```
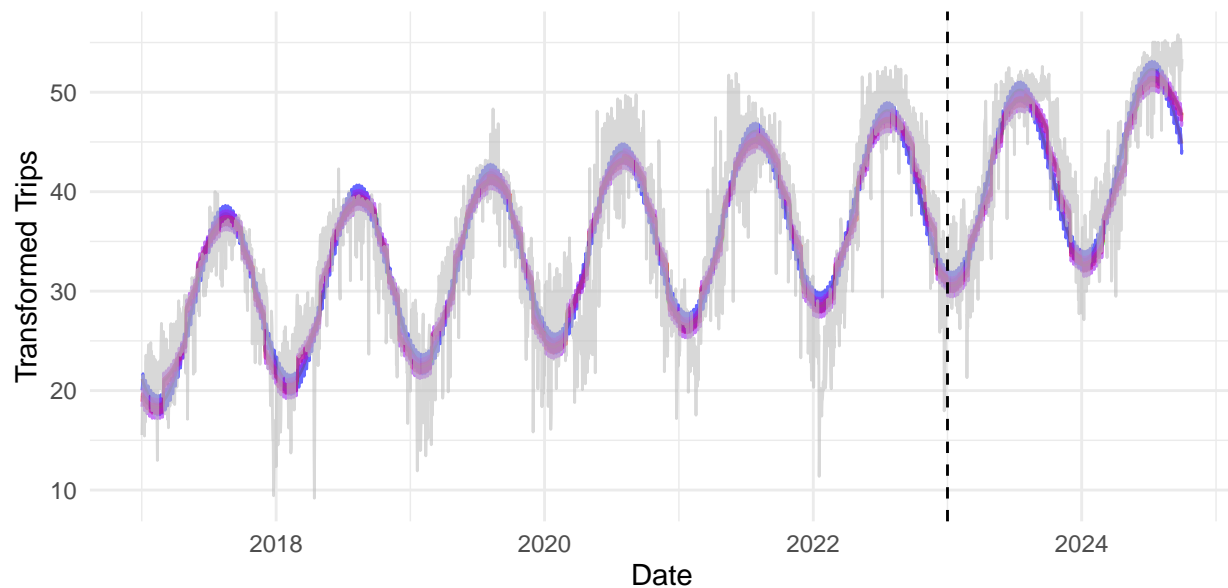
Model Fits vs. Actual Data (2017–2024)
## Models trained on data before 2023 (left of dashed line)


```r
# To see individual plots if the combined one is too cluttered:
# Function to plot one model against actuals
plot_one_model <- function(pred_col, model_name) {
  ggplot(plot_data, aes(x = trip_date)) +
    geom_line(aes(y = y_trans), color = "gray", alpha = 0.5) +
    geom_line(aes(y = .data[[pred_col]]), color = "blue") +
    geom_vline(xintercept = as.Date("2023-01-01"), linetype = "dashed") +
    labs(title = paste("Fit:", model_name), y = "Transformed Trips", x = "Date") +
    theme_minimal()
}

plot_one_model("Pred_Mod1", "Trend + Fourier")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.
```
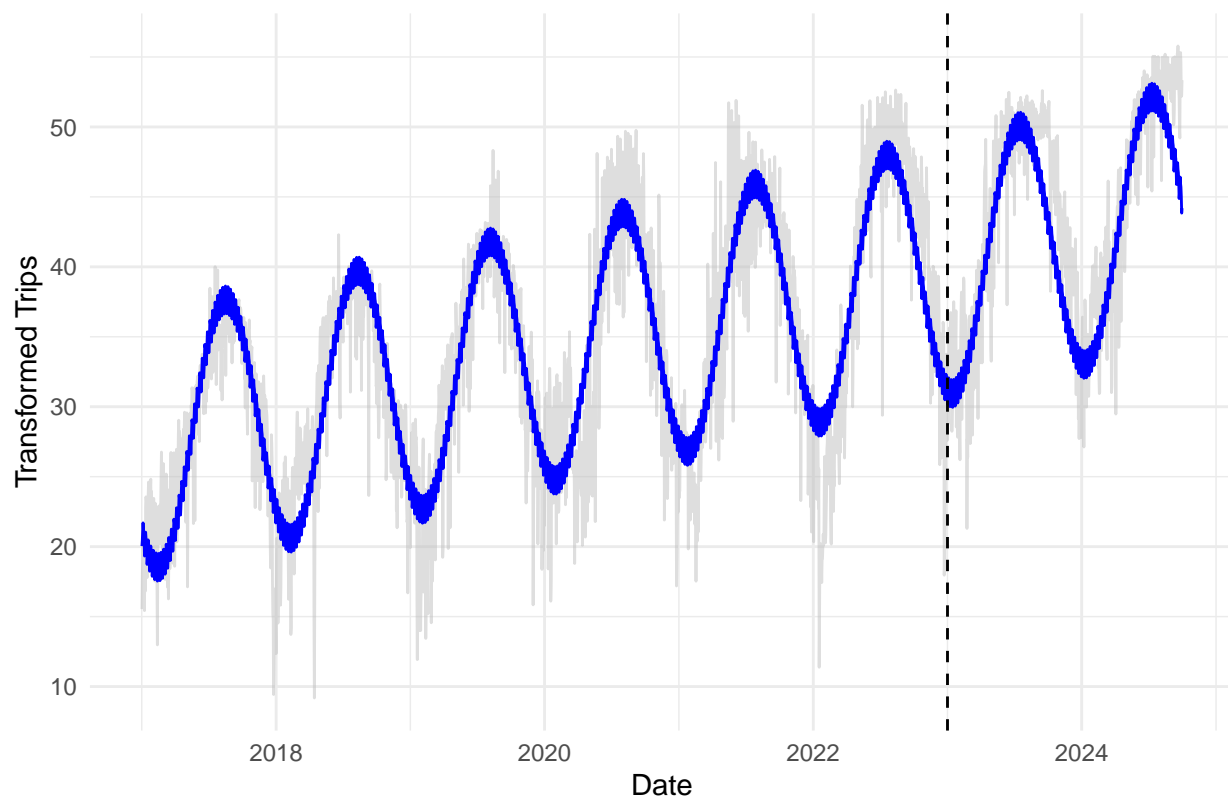
Fit: Trend + Fourier


```
plot_one_model("Pred_Mod2", "Trend + Fourier + DayOfWeek")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.
```
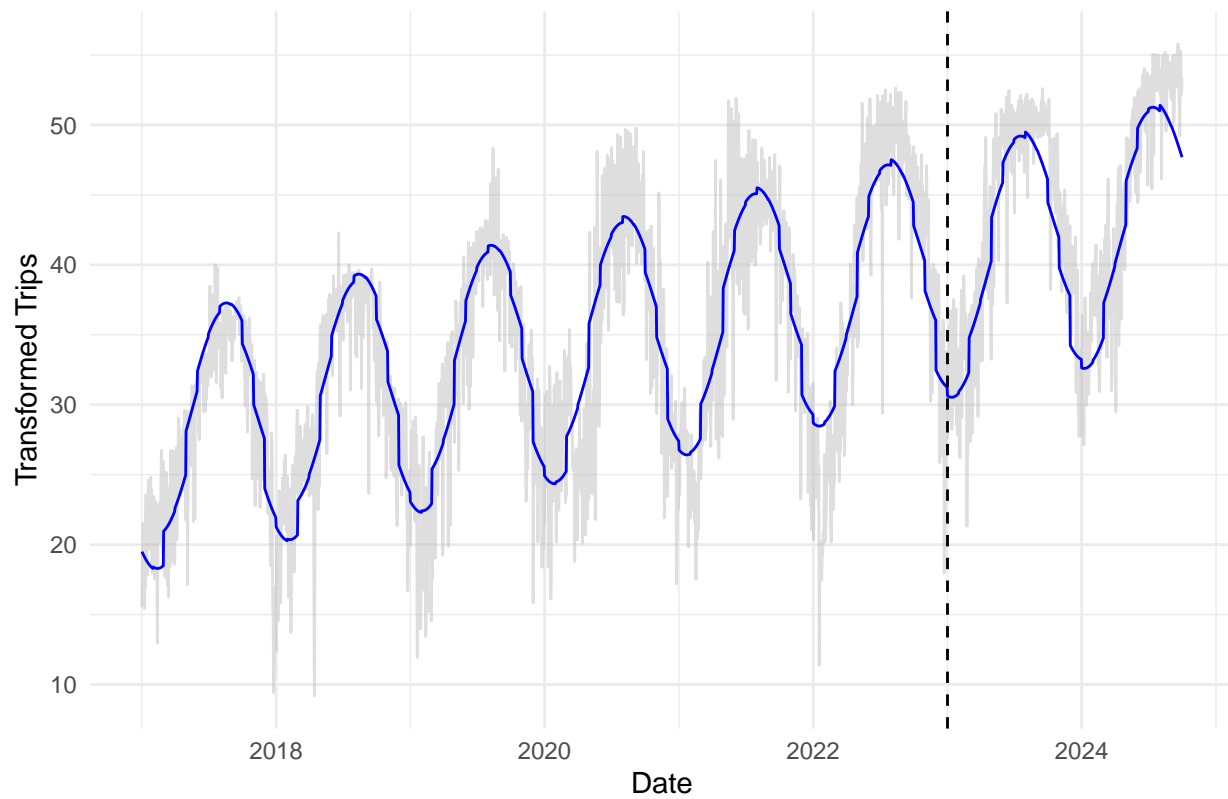
Fit: Trend + Fourier + DayOfWeek


```
plot_one_model("Pred_Mod3", "Trend + Fourier + Month")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.
```

Fit: Trend + Fourier + Month


```
plot_one_model("Pred_Mod4", "Trend + Fourier + Month + Weekend")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.
```
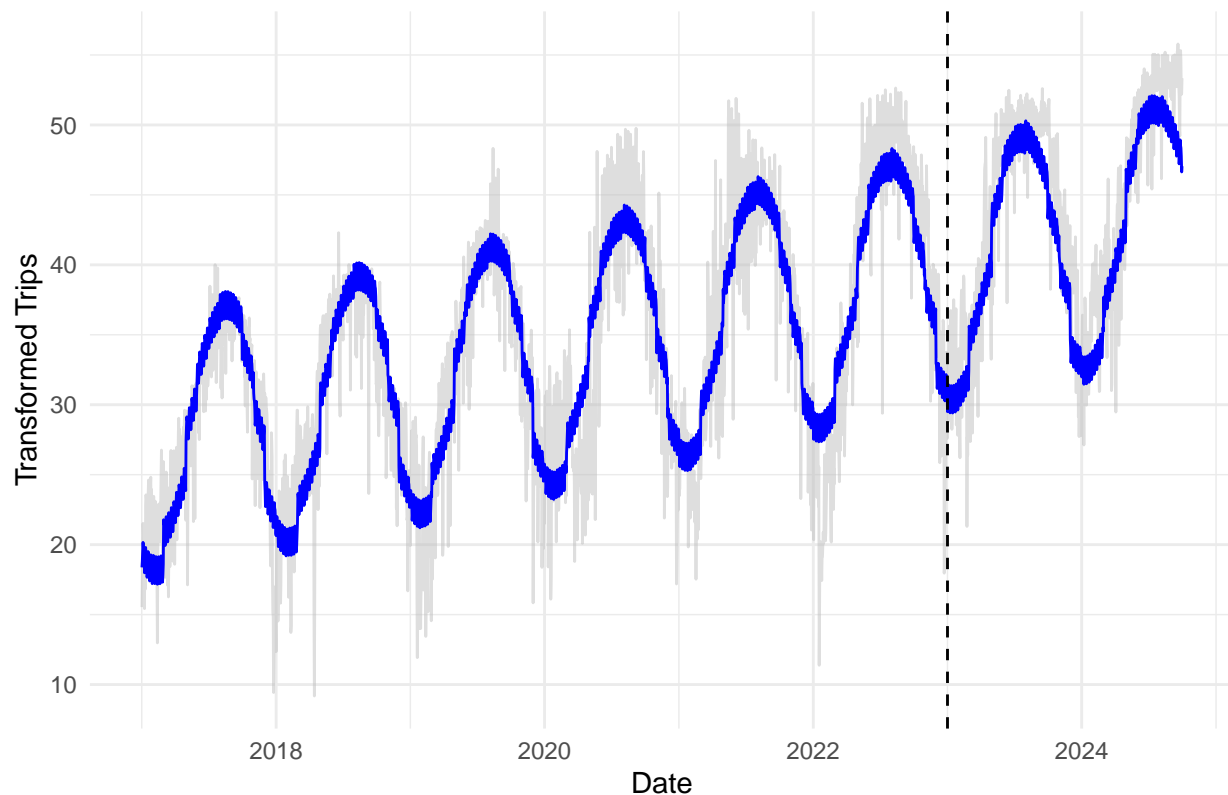
Fit: Trend + Fourier + Month + Weekend


```
plot_one_model("Pred_Mod5", "Trend + Fourier + Month + DayOfWeek")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.
```

Fit: Trend + Fourier + Month + DayOfWeek

Best Model Candidate is Trend + Fourier + Month + DayOfWeek.

Step 6: Select what degree polynomial for the trend

```r
# --- Step 6: Selecting Polynomial Trend Degree (Rolling Origin CV) ---

# Setup
max_degree <- 8
results_poly <- data.frame(Degree = 1:max_degree, APSE = NA)
min_train_size <- 730 # First 2 years as initial training window

# We assume the BEST structure found in Step 5 was:
# Trend + Fourier + Month + DayOfWeek
# So we fix those covariates and vary the polynomial degree 'p' for the Trend.

print("Running Rolling Origin CV for Trend with Best Seasonal Structure...")
```

```
## [1] "Running Rolling Origin CV for Trend with Best Seasonal Structure..."
```

```r
# Limit CV to pre-2024 data to avoid peeking at the final test set
cv_data <- subset(bike, trip_date < "2024-01-01")
n_cv <- nrow(cv_data)

for (p in 1:max_degree) {
  errors <- numeric()

  # Create polynomial basis for the FULL CV dataset first
  # raw=TRUE is safer for forecasting to avoid basis drift as n grows
  poly_basis <- poly(cv_data$time_index, p, raw = TRUE)
```

```r
  # Rolling Loop: Predict one day ahead, stepping by 30 days for speed
  for (i in seq(min_train_size, n_cv - 30, by = 30)) {

    # Define Train/Test indices
    train_idx <- 1:i
    test_idx <- (i + 1):(i + 30) # Testing a 30-day block

    # Build Dataframes manually
    # Incorporating: Poly Trend + Month + Weekday + Fourier (Sin/Cos)
    x_train <- data.frame(
      y = cv_data$y_trans[train_idx],
      poly = poly_basis[train_idx, , drop=FALSE],
      month = cv_data$month_fac[train_idx],
      weekday = cv_data$weekday_fac[train_idx],
      sin_year = cv_data$sin_year[train_idx]     )

    x_test <- data.frame(
      poly = poly_basis[test_idx, , drop=FALSE],
      month = cv_data$month_fac[test_idx],
      weekday = cv_data$weekday_fac[test_idx],
      sin_year = cv_data$sin_year[test_idx]     )

    # Fit Model
    # y ~ . uses all columns in x_train as predictors
    fit <- lm(y ~ ., data = x_train)

    # Predict
    preds <- predict(fit, newdata = x_test)

    # Calculate Squared Error for this window
    errors <- c(errors, (cv_data$y_trans[test_idx] - preds)^2)
  }

  # Store average error for degree 'p'
  results_poly$APSE[p] <- mean(errors)
  print(paste("Degree", p, "APSE:", round(results_poly$APSE[p], 4)))
}
```

```
## [1] "Degree 1 APSE: 18.7288"
## [1] "Degree 2 APSE: 18.6169"
## [1] "Degree 3 APSE: 19.6611"
## [1] "Degree 4 APSE: 21.3845"
## [1] "Degree 5 APSE: 24.4192"
## [1] "Degree 6 APSE: 28.9207"
## [1] "Degree 7 APSE: 36.3424"
## [1] "Degree 8 APSE: 44.0664"
```

```r
# Plot Results
plot(results_poly$Degree, results_poly$APSE, type='b', main="Polynomial Selection (with Best Seasonality
best_p <- which.min(results_poly$APSE)
points(best_p, results_poly$APSE[best_p], col = "red", cex = 2, pch = 1)
text(best_p, results_poly$APSE[best_p], paste("Best:", best_p), pos = 3, col = "red")
```

Polynomial Selection (with Best Seasonality)
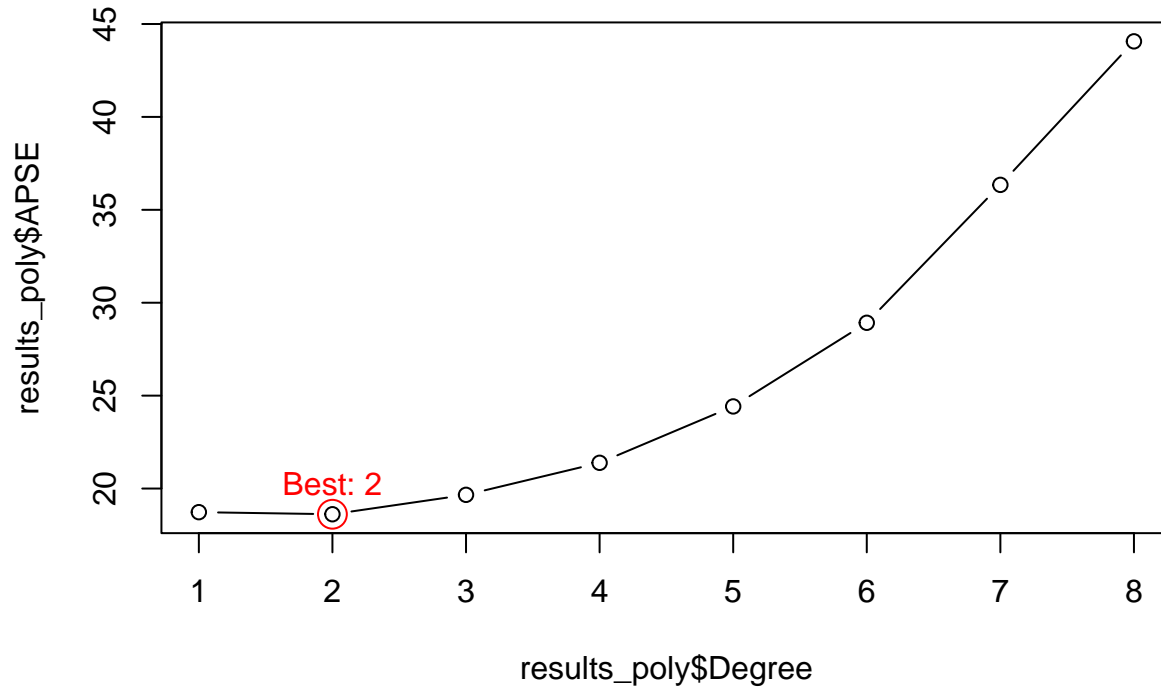

```r
print(paste("Best Polynomial Degree:", best_p))
```

```
## [1] "Best Polynomial Degree: 2"
```

As anticipated, the best model is a linear polynomial. We theorize that there might be some multicollinearity and use VIF to check:

```r
# --- Fit Final Regression Model (REQUIRED before Diagnostics) ---

# 1. Define the Training Data (2017-2023)
# We exclude the 2024 data (Test Set) to keep it unseen
train_full <- subset(bike, trip_date < "2024-01-01")

# 2. Create the polynomial basis for the trend
# We use Degree = 1 because your Cross-Validation identified it as the best.
poly_basis_final <- poly(train_full$time_index, 1, raw = TRUE)

# 3. Build the dataframe for lm
# We combine the trend with the seasonal components (Month + Weekday + Fourier)
final_train_data <- data.frame(
  y = train_full$y_trans,
  poly = poly_basis_final,
  month = train_full$month_fac,
  weekday = train_full$weekday_fac,
  sin_year = train_full$sin_year
)

# 4. Fit the model
# This creates the 'final_reg' object the next chunk is looking for
final_reg <- lm(y ~ ., data = final_train_data)
```

```r
# Check summary to confirm it worked
summary(final_reg)
```

```
##
## Call:
## lm(formula = y ~ ., data = final_train_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -18.4283  -1.9434   0.3893   2.4272  13.7409
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept) 20.9416583  0.5454663  38.392  < 2e-16 ***
## X1           0.0062725  0.0001039  60.385  < 2e-16 ***
## month2      -0.0214172  0.3774770  -0.057 0.954759
## month3       2.2110517  0.3832445   5.769 8.93e-09 ***
## month4       3.1891624  0.4813278   6.626 4.20e-11 ***
## month5       6.5118087  0.6539711   9.957  < 2e-16 ***
## month6       8.0104879  0.8313682   9.635  < 2e-16 ***
## month7       8.4153838  0.9351974   8.999  < 2e-16 ***
## month8       8.8832011  0.9313077   9.538  < 2e-16 ***
## month9       9.0271892  0.8212392  10.992  < 2e-16 ***
## month10      7.1104829  0.6416023  11.082  < 2e-16 ***
## month11      4.5313766  0.4726663   9.587  < 2e-16 ***
## month12      0.6219054  0.3819394   1.628 0.103589
## weekdayMon   0.4996046  0.2837193   1.761 0.078374 .
## weekdayTue   1.3681208  0.2837276   4.822 1.51e-06 ***
## weekdayWed   1.9604860  0.2837372   6.910 6.13e-12 ***
## weekdayThu   1.6845539  0.2837378   5.937 3.30e-09 ***
## weekdayFri   1.5369271  0.2837273   5.417 6.63e-08 ***
## weekdaySat   1.0011027  0.2837199   3.528 0.000425 ***
## sin_year    -4.3837879  0.4579141  -9.573  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.835 on 2536 degrees of freedom
## Multiple R-squared:  0.8198, Adjusted R-squared:  0.8184
## F-statistic: 607.2 on 19 and 2536 DF,  p-value: < 2.2e-16
```

```r
# --- Multicollinearity Check (VIF) ---
library(car)

# Calculate VIF for the final regression model
# Note: We use generalized VIF (GVIF) because 'month' and 'weekday' are categorical factors
vif_values <- vif(final_reg)

print(vif_values)
```

```
##               GVIF Df GVIF^(1/(2*Df))
## X1         1.020741  1        1.010317
## month     18.691478 11        1.142357
## weekday    1.000795  6        1.000066
## sin_year  18.343028  1        4.282876
```

17

```r
# Check for high multicollinearity
# Rule of thumb: GVIF^(1/(2*Df)) > 2 indicates concerning collinearity (equivalent to VIF > 4)
high_vif <- vif_values[,"GVIF^(1/(2*Df))"] > 2
print("Variables with potentially problematic collinearity:")
```

```
## [1] "Variables with potentially problematic collinearity:"
```

```r
print(vif_values[high_vif, ])
```

```
##            GVIF           Df GVIF^(1/(2*Df))
##       18.343028     1.000000       4.282876
```

```r
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.3.3
```

```
## Loading required package: Matrix
```

```
##
## Attaching package: 'Matrix'
```

```
## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack
```

```
## Loaded glmnet 4.1-9
```

```r
# 1. Build the EXACT same design matrix as the OLS model
X_train <- model.matrix(
  y_trans ~ poly(time_index, 1, raw = TRUE) + month_fac + weekday_fac + sin_year,
  data = train_full
)[, -1]   # remove intercept (glmnet adds its own)

Y_train <- train_full$y_trans

# 2. Fit CV-Lasso on the exact same feature matrix
set.seed(443)
cv_lasso <- cv.glmnet(
  X_train,
  Y_train,
  alpha = 1,
  standardize = TRUE,
  intercept = TRUE,
  family = "gaussian"
)

best_lambda <- cv_lasso$lambda.min
best_lambda
```

```
## [1] 0.002414156
```

```r
lasso_model <- glmnet(
  X_train, Y_train,
  alpha = 1,
  lambda = best_lambda,
  standardize = TRUE,
  intercept = TRUE,
  family = "gaussian"
```

```
)

coef(lasso_model)

## 20 x 1 sparse Matrix of class "dgCMatrix"
##                                       s0
## (Intercept)                  21.229863999
## poly(time_index, 1, raw = TRUE)  0.006271939
## month_fac2                   -0.088734213
## month_fac3                    2.083074848
## month_fac4                    2.979351032
## month_fac5                    6.207111093
## month_fac6                    7.623008014
## month_fac7                    7.980442541
## month_fac8                    8.450384992
## month_fac9                    8.644152951
## month_fac10                   6.812015566
## month_fac11                   4.328646017
## month_fac12                   0.498842295
## weekday_facMon                0.456425154
## weekday_facTue                1.324743590
## weekday_facWed                1.917580696
## weekday_facThu                1.642094917
## weekday_facFri                1.494253243
## weekday_facSat                0.959006696
## sin_year                     -4.572277428
```

We applied Lasso regression to test for feature redundancy. The optimal lambda was small (0.0045), and Lasso retained nearly all predictors, confirming their relevance. The validation RMSE for Lasso (3.904) was virtually identical to the OLS model (3.903). Since regularization yielded no performance gain or significant simplification, we proceeded with the standard OLS regression model for the final forecasting phase.

```r
# --- Consistent AIC/BIC for OLS vs Lasso using Gaussian likelihood ---

# Unified AIC / BIC functions ------------------------------------------

calc_AIC <- function(rss, n, k) {
  sigma2 <- rss / n
  loglik <- -n/2 * (log(2*pi*sigma2) + 1)
  AIC <- -2*loglik + 2*k
  return(AIC)
}

calc_BIC <- function(rss, n, k) {
  sigma2 <- rss / n
  loglik <- -n/2 * (log(2*pi*sigma2) + 1)
  BIC <- -2*loglik + log(n)*k
  return(BIC)
}

calc_AICc <- function(aic, n, k) {
  aic + (2*k*(k+1)) / (n - k - 1)
}

# ----------------------------------------------------------------------
```

```r
# 1. OLS model (final_reg)
# -----------------------------------------------------------------

rss_ols <- sum(residuals(final_reg)^2)
n_ols <- nobs(final_reg)
k_ols <- length(coef(final_reg))

AIC_ols  <- calc_AIC(rss_ols, n_ols, k_ols)
BIC_ols  <- calc_BIC(rss_ols, n_ols, k_ols)
AICc_ols <- calc_AICc(AIC_ols, n_ols, k_ols)
adjR2_ols <- summary(final_reg)$adj.r.squared


# -----------------------------------------------------------------
# 2. Lasso model
# -----------------------------------------------------------------

# training predictions
lasso_pred_train <- predict(lasso_model, newx = X_train)

rss_lasso <- sum((Y_train - lasso_pred_train)^2)
n_lasso <- length(Y_train)
k_lasso <- sum(coef(lasso_model) != 0)  # df = non-zero coefficients

AIC_lasso  <- calc_AIC(rss_lasso, n_lasso, k_lasso)
BIC_lasso  <- calc_BIC(rss_lasso, n_lasso, k_lasso)
AICc_lasso <- calc_AICc(AIC_lasso, n_lasso, k_lasso)

# Adjusted R2
tss <- sum((Y_train - mean(Y_train))^2)
adjR2_lasso <- 1 - ((rss_lasso/(n_lasso - k_lasso - 1)) /
                    (tss/(n_lasso - 1)))

# -----------------------------------------------------------------
# 3. Final comparison table
# -----------------------------------------------------------------

comparison <- data.frame(
  Model = c("OLS", "Lasso"),
  AIC = c(AIC_ols, AIC_lasso),
  AICc = c(AICc_ols, AICc_lasso),
  BIC = c(BIC_ols, BIC_lasso),
  Adjusted_R2 = c(adjR2_ols, adjR2_lasso)
)

comparison
```

```
##   Model      AIC     AICc      BIC Adjusted_R2
## 1   OLS 14145.34 14145.67 14262.27   0.8184397
## 2 Lasso 14145.65 14145.98 14262.57   0.8183464
```

Step 7: Residual Diagnostics on the Selected Model

```r
# --- 1. Define Class Diagnostic Functions ---

RegressionDiagnosicsPlots <- function(model, my_residuals=NULL, my_fitted=NULL, title_suffix="") {
```

```r
  # Handle both standard lm models and manually passed residuals (for flexibility)
  if(is.null(my_residuals)) {
    my_residuals <- residuals(model)
    my_fitted <- fitted(model)
    model_name <- deparse(substitute(model))
  } else {
    model_name <- title_suffix
  }

  # 1. Histogram
  hist(my_residuals, xlab = "Residuals", main = paste("Hist:", model_name))

  # 2. QQ Plot
  car::qqPlot(my_residuals, pch = 16, col = adjustcolor("black", 0.7),
              xlab = "Theoretical Quantiles (Normal)", ylab = "Sample Quantiles",
              main = "Normal Q-Q Plot")

  # 3. Fitted vs Residuals
  if(!is.null(my_fitted)){
    plot(my_fitted, my_residuals, pch = 16, col = adjustcolor("black", 0.5),
         xlab = "Fitted Values", ylab = "Residuals", main = "Fitted vs Residuals")
    abline(h = 0, lty = 2, col = 'red')
  } else {
    plot.new() # Empty plot if no fitted values
  }

  # 4. Time vs Residuals
  plot(my_residuals, pch = 16, col = adjustcolor("black", 0.5),
       xlab = "Time", ylab = "Residuals", main = "Residuals vs Time")
  abline(h = 0, lty = 2, col = 'red')

  # 5. ACF
  acf(my_residuals, main = "ACF of Residuals")

  par(mfrow = c(1, 1)) # Reset layout
}

RegressionDiagnosicsTests <- function(model, my_residuals=NULL, segments=6) {

  if(is.null(my_residuals)) {
    my_residuals <- residuals(model)
  }

  print("--- Shapiro-Wilk Test (Normality) ---")
  # Shapiro test limit is 5000
  if(length(my_residuals) > 5000) {
    print(shapiro.test(sample(my_residuals, 5000)))
  } else {
    print(shapiro.test(my_residuals))
  }

  print("--- Kolmogorov-Smirnov Test (Normality) ---")
  print(ks.test(scale(my_residuals), "pnorm"))
```

```
  print("--- Fligner-Killeen Test (Homoscedasticity) ---")
  # Create segments dynamically based on data length
  n <- length(my_residuals)
  # Make segments roughly equal size
  seg <- factor(cut(1:n, breaks = segments, labels = FALSE))
  print(fligner.test(my_residuals, seg))

  print("--- Runs Test (Randomness) ---")
  par(mfrow = c(1, 1))
  print(randtests::runs.test(my_residuals))
}



# --- 2. Apply to Your Final Regression Model ---

# Assuming 'final_reg' is your best model from the previous steps
# (Trend + Fourier + Month + DayOfWeek)

# Generate Plots
RegressionDiagnosicsPlots(final_reg)
```


Hist: final_reg

Normal Q–Q Plot

Fitted vs Residuals

Residuals vs Time

ACF of Residuals


```
# Generate Formal Tests
RegressionDiagnosicsTests(final_reg)
```

```
## [1] "--- Shapiro-Wilk Test (Normality) ---"
##
##  Shapiro-Wilk normality test
```

```
## 
## data:  my_residuals
## W = 0.97393, p-value < 2.2e-16
## 
## [1] "--- Kolmogorov-Smirnov Test (Normality) ---"
## 
##  Asymptotic one-sample Kolmogorov-Smirnov test
## 
## data:  scale(my_residuals)
## D = 0.064303, p-value = 1.322e-09
## alternative hypothesis: two-sided
## 
## [1] "--- Fligner-Killeen Test (Homoscedasticity) ---"
## 
##  Fligner-Killeen test of homogeneity of variances
## 
## data:  my_residuals and seg
## Fligner-Killeen:med chi-squared = 25.293, df = 5, p-value = 0.0001223
## 
## [1] "--- Runs Test (Randomness) ---"
## 
##  Runs Test
## 
## data:  my_residuals
## statistic = -21.683, runs = 731, n1 = 1278, n2 = 1278, n = 2556,
## p-value < 2.2e-16
## alternative hypothesis: nonrandomness
```

Diagnostic tests confirm that while the regression model explains a large portion of the variance, the residuals significantly violate assumptions of normality, constant variance, and independence. This lack of independence specifically motivates the use of ARIMA modeling in the next phase to capture the remaining autocorrelation structure.

```r
# Residual Diagnostics for Lasso

# Fitted values for training data
fitted_lasso <- predict(lasso_model, newx = X_train)

# Residuals
resid_lasso <- Y_train - fitted_lasso

plot(fitted_lasso, resid_lasso,
     xlab = "Fitted Values", ylab = "Residuals",
     main = "Lasso: Residuals vs Fitted")
abline(h = 0, col = "red")
```

Lasso: Residuals vs Fitted


```
qqnorm(resid_lasso)
qqline(resid_lasso, col = "red")
```
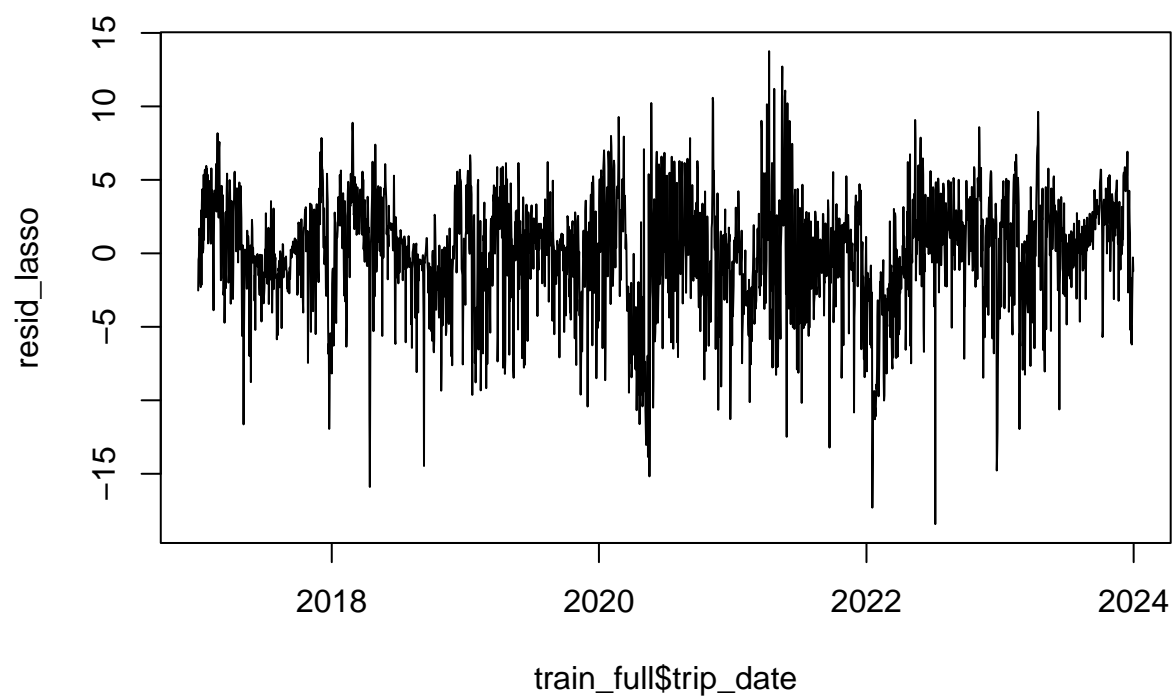
Normal Q–Q Plot


```
hist(resid_lasso, breaks = 30)
```

26

Histogram of resid_lasso


```r
plot(train_full$trip_date, resid_lasso, type="l")
```



```r
acf(resid_lasso)
```

s0