

Step 1: Loading in the Data

```
# Packages used in tutorials
library(MASS)      # boxcox
library(car)        # qqPlot
library(randtests) # runs.test
# library(forecast) # OPTIONAL if you want auto.arima, not required

bike <- read.csv("trips_per_day.csv")
bike <- bike %>%
  filter(!is.na(trip_date))

bike$trip_date <- as.Date(bike$trip_date)

# Filter the bike dataframe to keep only trips from 2017 onwards
bike <- bike[bike$trip_date >= as.Date("2017-01-01"), ]

# Verify the new range
range(bike$trip_date)

## [1] "2017-01-01" "2024-09-30"

str(bike)

## 'data.frame': 2830 obs. of 2 variables:
## $ trip_date: Date, format: "2017-01-01" "2017-01-02" ...
## $ n_trips  : int 487 775 918 1359 1202 1189 632 477 1269 622 ...
head(bike)

##    trip_date n_trips
## 1 2017-01-01     487
## 2 2017-01-02     775
## 3 2017-01-03     918
## 4 2017-01-04    1359
## 5 2017-01-05    1202
## 6 2017-01-06    1189

range(bike$trip_date)

## [1] "2017-01-01" "2024-09-30"

Initial Plotting for Time Series

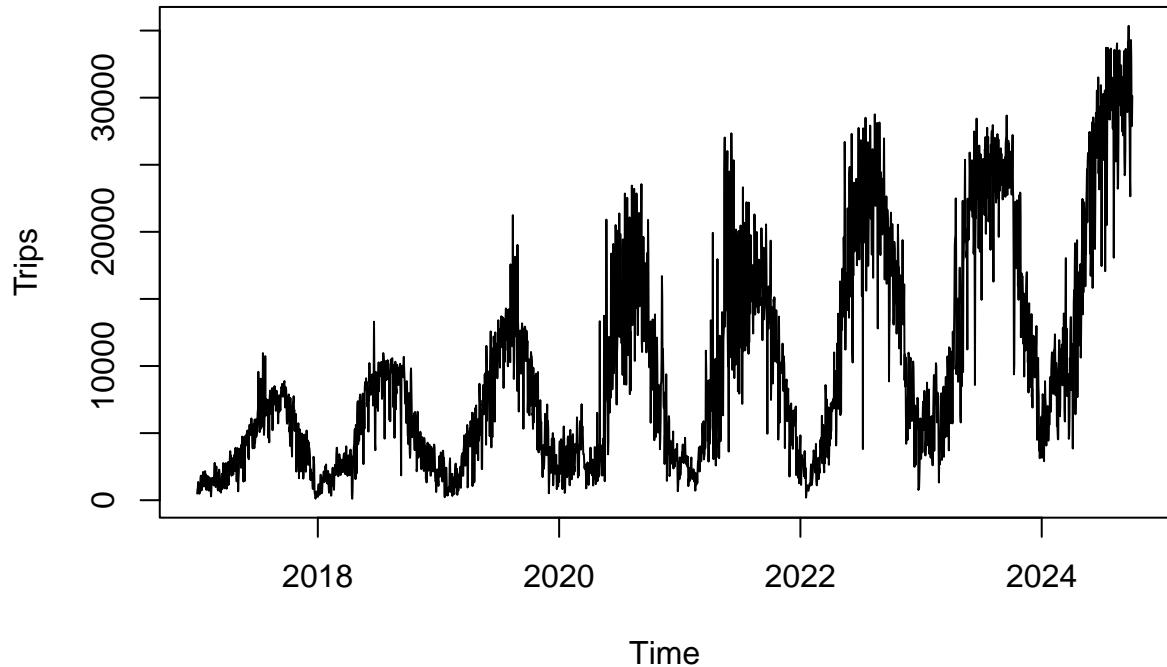
# Sort just in case
bike <- bike[order(bike$trip_date), ]

# Extract response as a vector
y <- bike$n_trips

# Daily frequency with yearly seasonality (approx 365)
bike_ts <- ts(
  y,
  start = c(as.numeric(format(min(bike$trip_date), "%Y")),
            as.numeric(format(min(bike$trip_date), "%j"))),
  frequency = 365
)
```

```
plot(bike_ts, main = "Daily BikeShare Trips in Toronto", ylab = "Trips")
```

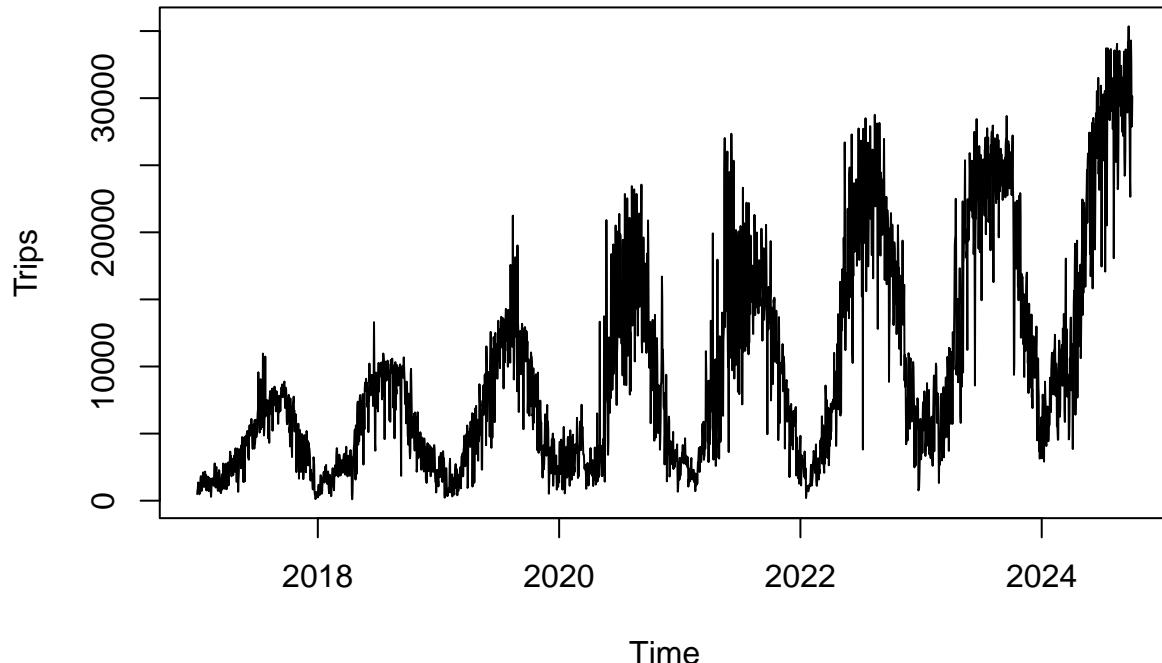
Daily BikeShare Trips in Toronto



Step 2: EDA

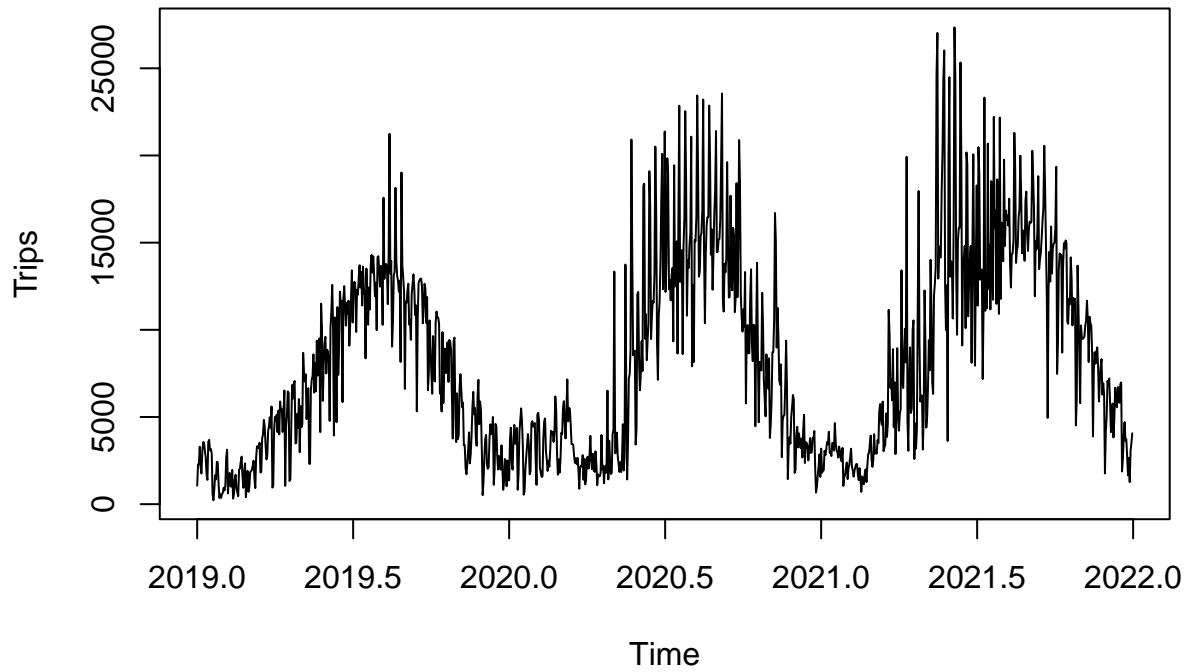
```
par(mfrow = c(1, 1))
plot(bike_ts, main = "Daily Trips", ylab = "Trips")
```

Daily Trips



```
# maybe a zoom on a couple of years
plot(window(bike_ts, start = c(2019, 1), end = c(2021, 365)),
     main = "Daily Trips: 2019–2021", ylab = "Trips")
```

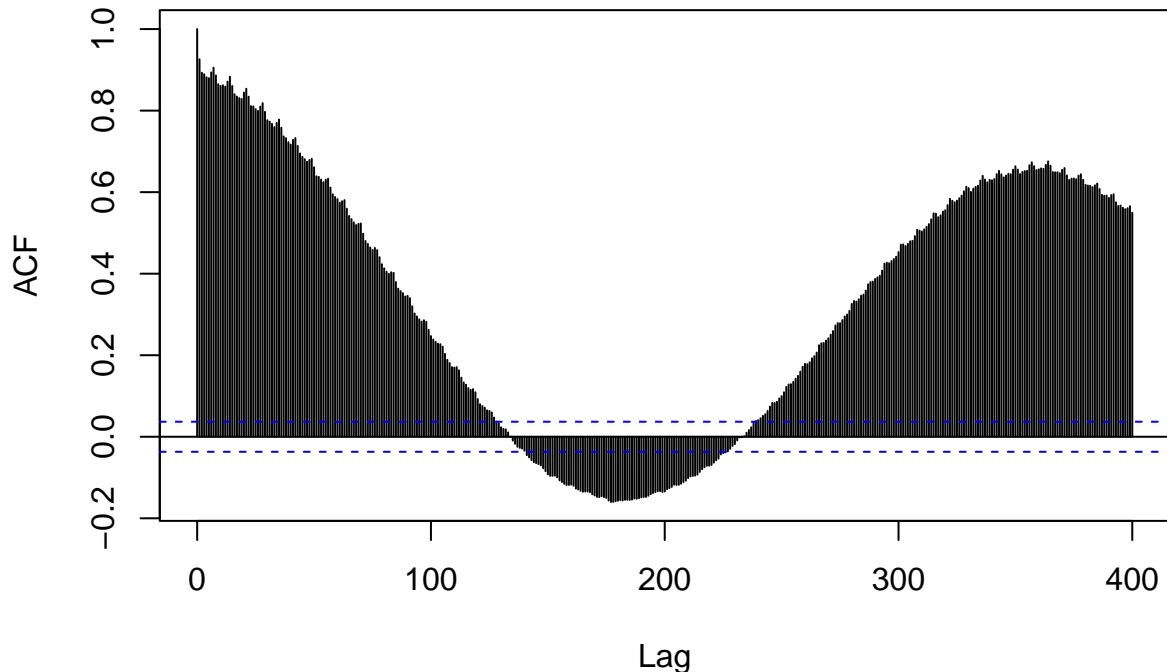
Daily Trips: 2019–2021



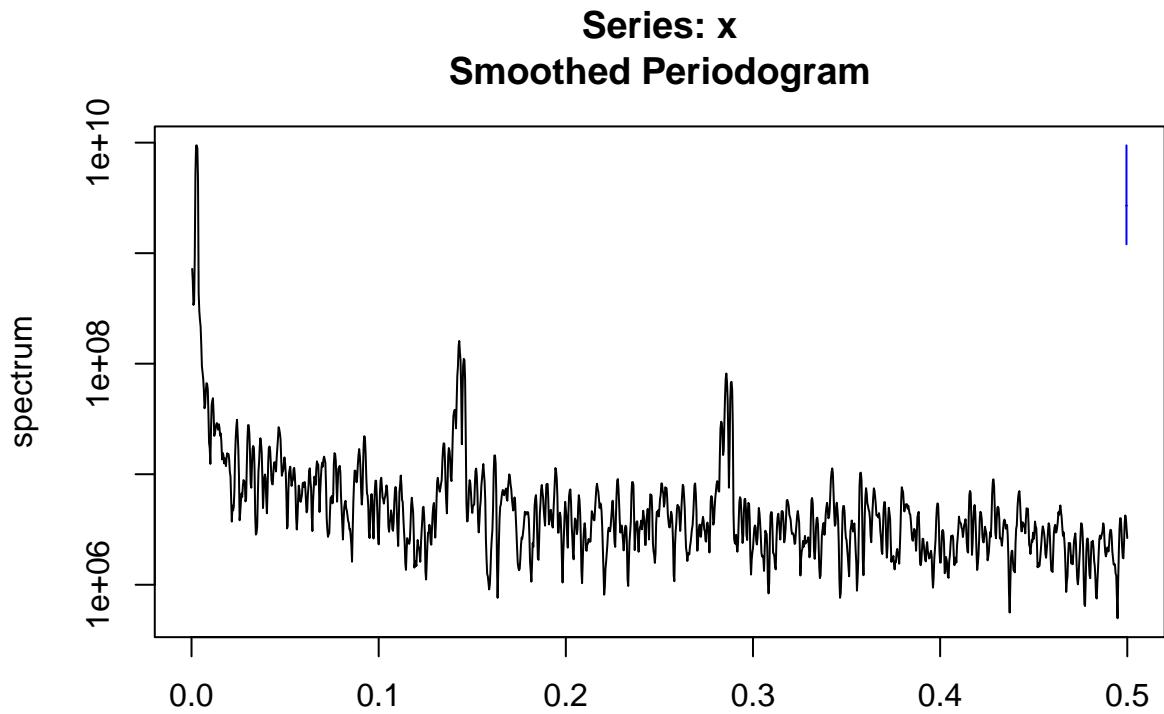
Confirming Seasonality with ACF and Spectrum

```
acf(as.vector(bike_ts), lag.max = 400,  
    main = "ACF of Daily Trips")
```

ACF of Daily Trips



```
spec_bike <- spectrum(as.vector(bike_ts), spans = 5)
```

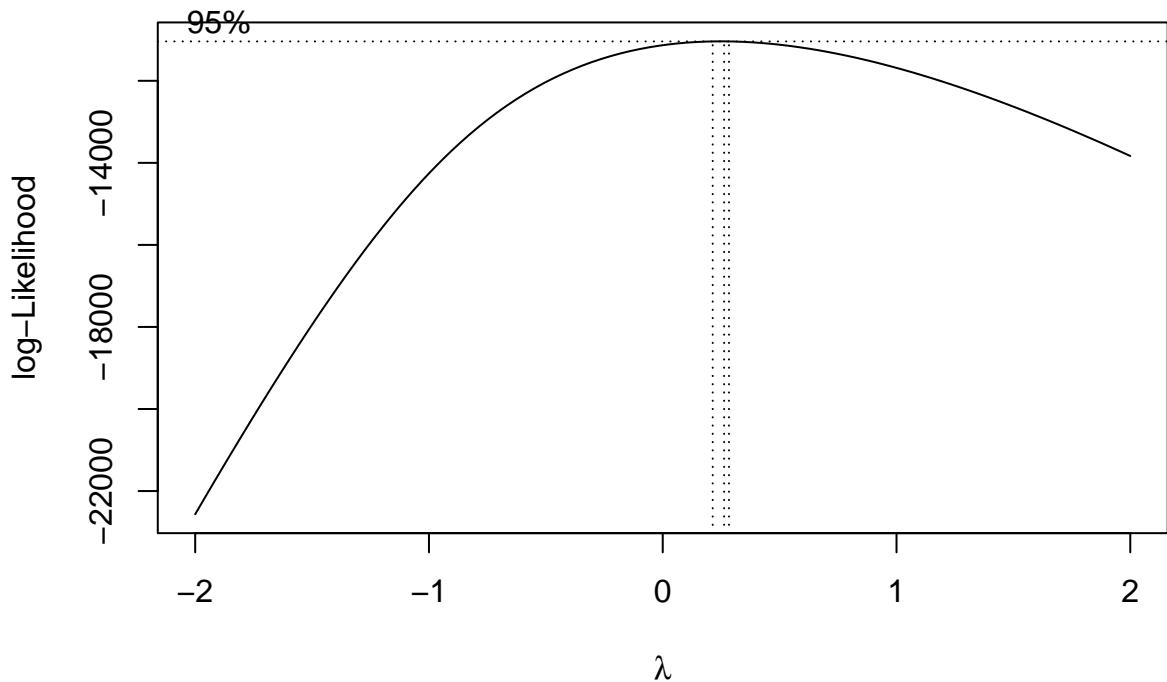


```
1 / spec_bike$freq[which.max(spec_bike$spec)] # estimated period
```

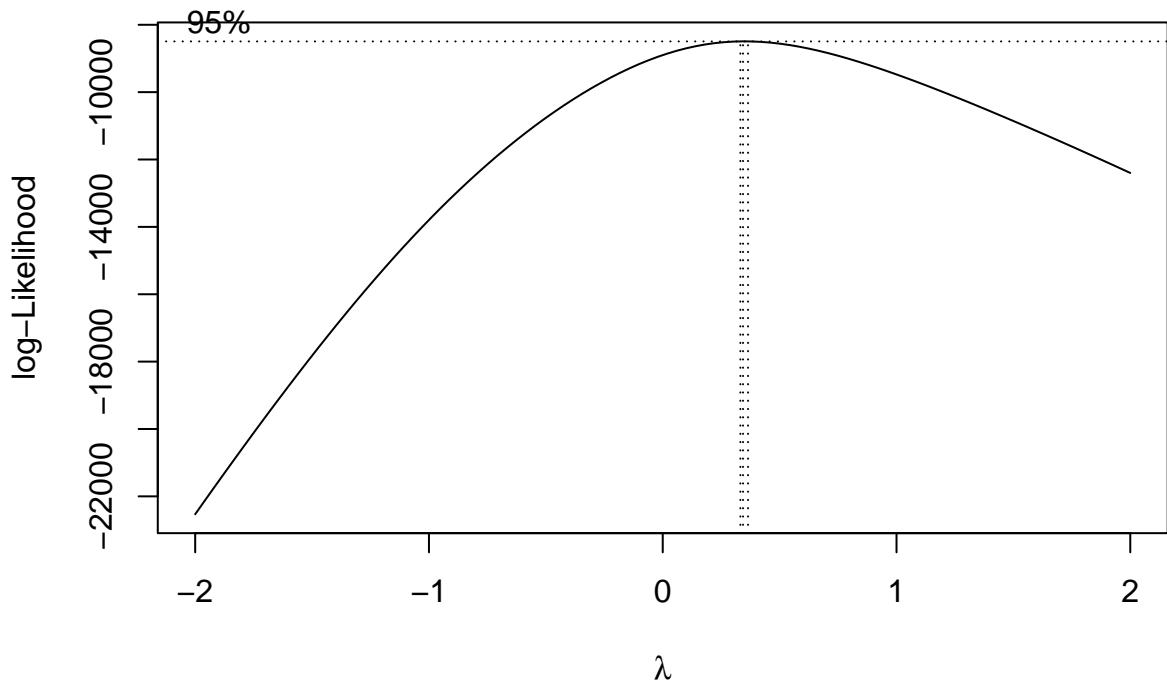
```
## [1] 360
```

Step 3: Box-Cox transformation

```
# Simple intercept-only model (like in tutorial)
bc_model_raw <- lm(bike_ts ~ 1)
boxcox_raw <- MASS::boxcox(bc_model_raw, lambda = seq(-2, 2, 0.1))
```

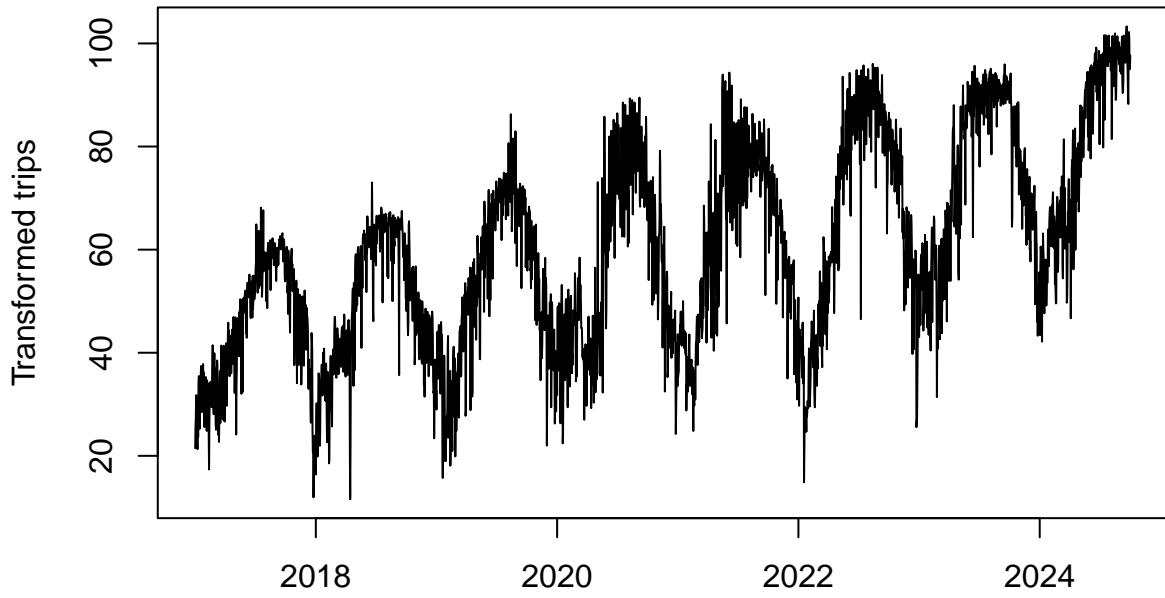


```
(lambda_opt_raw <- boxcox_raw$x[which.max(boxcox_raw$y)])  
  
## [1] 0.2626263  
tim <- time(bike_ts) # continuous time index  
  
# Season: year and day-of-year or month; simplest is month  
# Build a monthly factor from dates (instead of cycle, since this is daily)  
month <- factor(format(bike$trip_date, "%m"))  
  
reg_for_bc <- lm(bike_ts ~ tim + month)  
boxcox_mod <- MASS::boxcox(reg_for_bc, lambda = seq(-2, 2, 0.1))
```



```
(lambda_opt_mod <- boxcox_mod$x[which.max(boxcox_mod$y)])  
  
## [1] 0.3434343  
lam <- lambda_opt_mod # keep this for later  
  
if (lam == 0) {  
  y_trans <- log(bike_ts)  
} else if (lam > 0) {  
  y_trans <- (bike_ts^lam - 1) / lam  
} else {  
  # negative lambda + use minus sign trick like in lectures  
  y_trans <- -(bike_ts^lam)  
}  
  
plot(y_trans, main = "Transformed Daily Trips", ylab = "Transformed trips")
```

Transformed Daily Trips



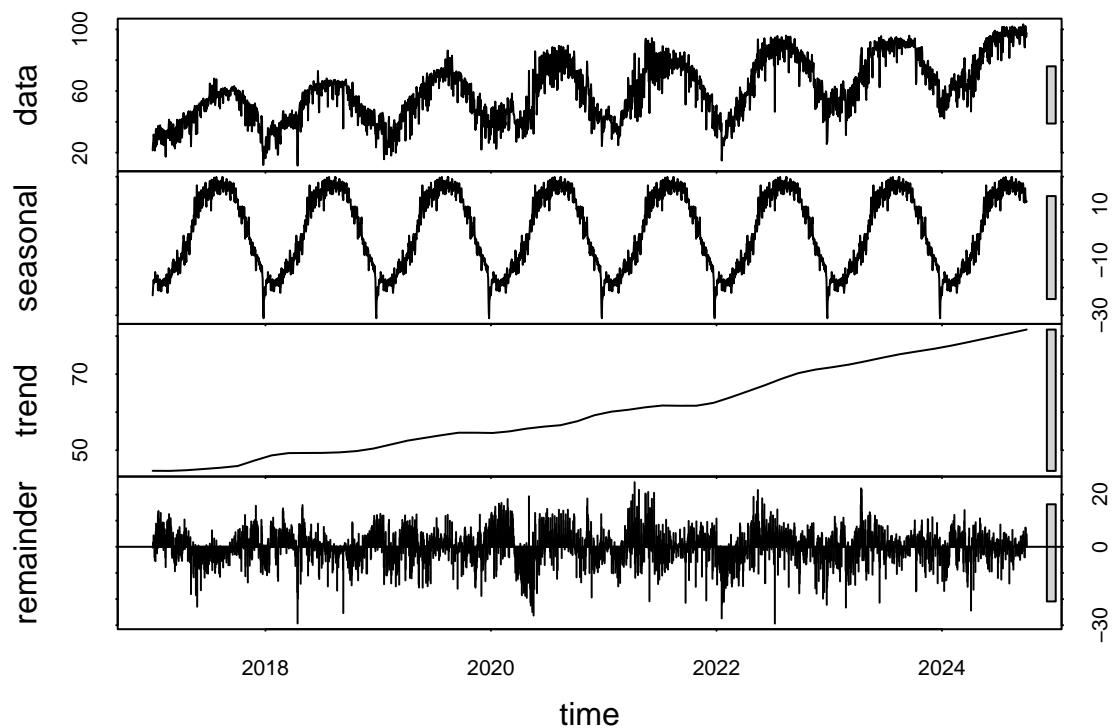
cal Decomposition

```
Decomp_bike <- stl(y_trans, s.window = "periodic")
plot(Decomp_bike, main = "STL Decomposition of Transformed Daily Trips")
```

Time

Classi-

STL Decomposition of Transformed Daily Trips



```

# Extract components if needed
bike_seasonal <- Decomp_bike$time.series[, "seasonal"]
bike_trend     <- Decomp_bike$time.series[, "trend"]
bike_remainder <- Decomp_bike$time.series[, "remainder"]

# --- Feature Engineering ---
# Time Index
tim <- time(bike_ts)

# Seasonal Factors
month <- factor(format(bike$trip_date, "%m"))
weekday_num <- as.numeric(format(bike$trip_date, "%u")) # 1=Mon, ..., 7=Sun
weekday <- factor(weekday_num)
is_weekend <- factor(ifelse(weekday_num >= 6, "Weekend", "Weekday"))

# Dataframe
model_data <- data.frame(
  y_trans = as.numeric(y_trans),
  tim = as.numeric(tim),
  month = month,
  weekday = weekday,
  is_weekend = is_weekend
)

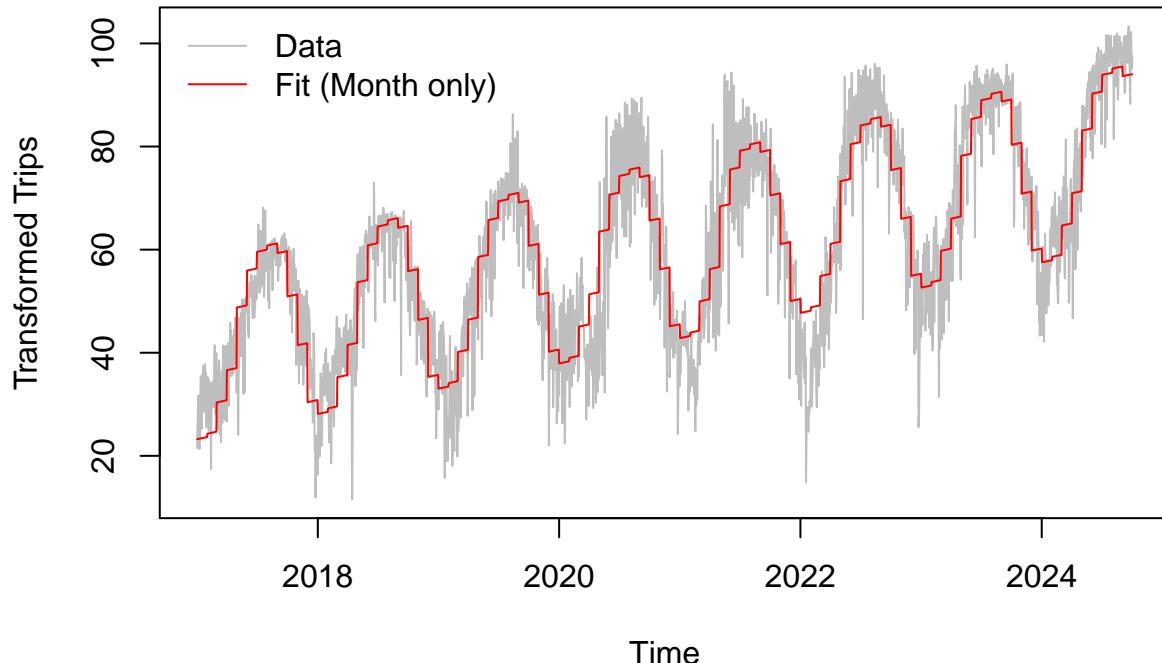
# --- Visualizing the Seasonality Tests ---

# Model 0: Base (Trend + Month Only)
mod_base <- lm(y_trans ~ tim + month, data = model_data)

# PLOT 0: Base Model
plot(model_data$tim, model_data$y_trans, type = "l", col = "gray",
      main = "Model 0: Monthly Seasonality Only", ylab = "Transformed Trips", xlab = "Time")
lines(model_data$tim, fitted(mod_base), col = "red", lwd = 1)
legend("topleft", legend = c("Data", "Fit (Month only)"), col = c("gray", "red"), lty = 1, bty = "n")

```

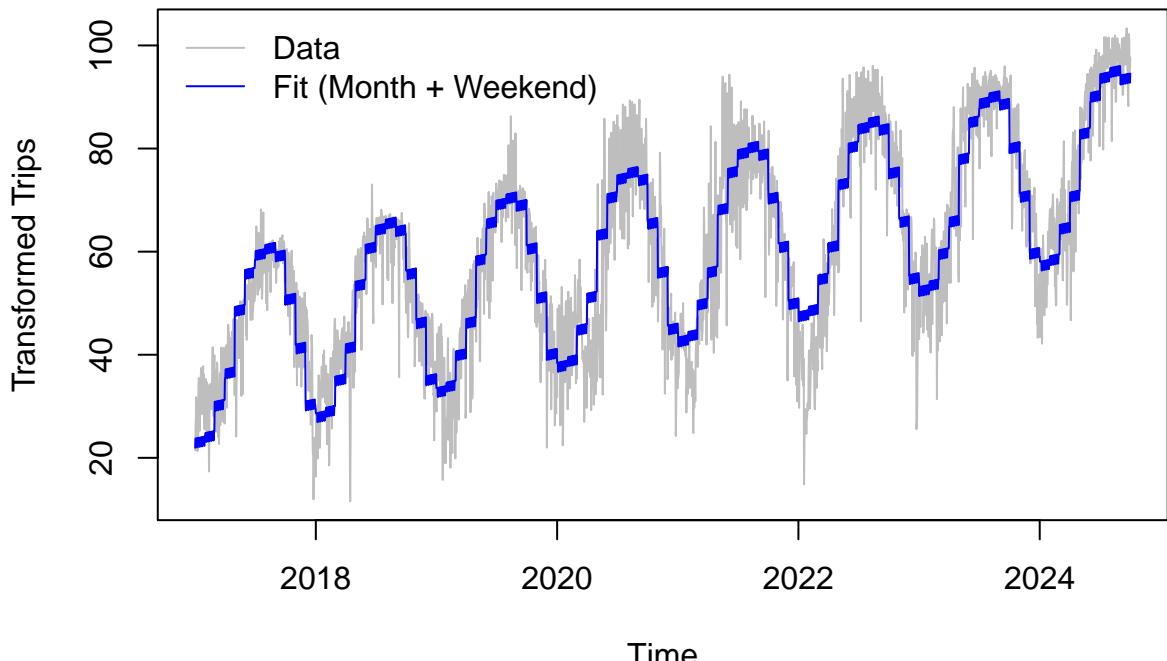
Model 0: Monthly Seasonality Only



```
# Model 1: Weekend (Trend + Month + Weekend Indicator)
mod_weekend <- lm(y_trans ~ tim + month + is_weekend, data = model_data)

# PLOT 1: Weekend Model
plot(model_data$tim, model_data$y_trans, type = "l", col = "gray",
      main = "Model 1: Month + Weekend Indicator", ylab = "Transformed Trips", xlab = "Time")
lines(model_data$tim, fitted(mod_weekend), col = "blue", lwd = 1)
legend("topleft", legend = c("Data", "Fit (Month + Weekend)"), col = c("gray", "blue"), lty = 1, bty = "n")
```

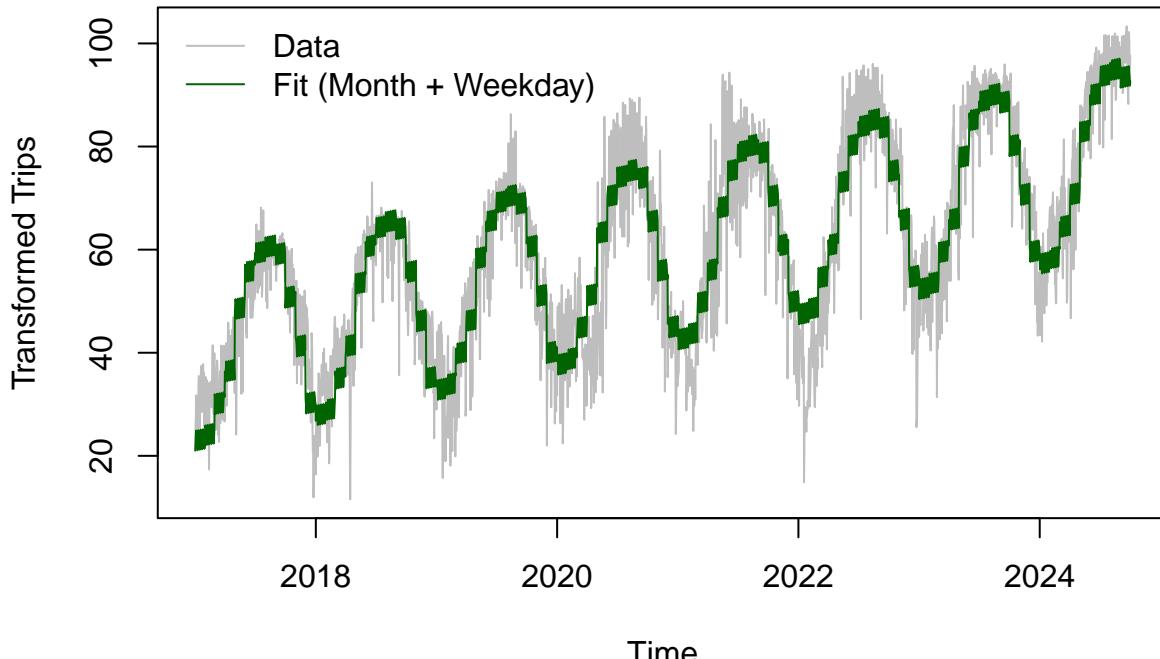
Model 1: Month + Weekend Indicator



```
# Model 2: Full Weekday (Trend + Month + Full Weekday)
mod_weekday <- lm(y_trans ~ tim + month + weekday, data = model_data)

# PLOT 2: Weekday Model
# We zoom in on a shorter window (e.g., first year) for this plot
# because the daily wiggle is hard to see on the full plot
plot(model_data$tim, model_data$y_trans, type = "l", col = "gray",
      main = "Model 2: Month + Full Weekday (First Year Zoom)", ylab = "Transformed Trips", xlab = "Time")
lines(model_data$tim, fitted(mod_weekday), col = "darkgreen", lwd = 1)
legend("topleft", legend = c("Data", "Fit (Month + Weekday)"), col = c("gray", "darkgreen"), lty = 1, b
```

Model 2: Month + Full Weekday (First Year Zoom)



```
# --- Hypothesis Testing ---
# Now you perform the tests to confirm what you see in the plots
anova(mod_base, mod_weekend)
```

```
## Analysis of Variance Table
##
## Model 1: y_trans ~ tim + month
## Model 2: y_trans ~ tim + month + is_weekend
##   Res.Df   RSS Df Sum of Sq    F    Pr(>F)
## 1    2817 172833
## 2    2816 171030  1     1803.5 29.694 5.496e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
anova(mod_weekend, mod_weekday)
```

```
## Analysis of Variance Table
##
## Model 1: y_trans ~ tim + month + is_weekend
## Model 2: y_trans ~ tim + month + weekday
##   Res.Df   RSS Df Sum of Sq    F    Pr(>F)
## 1    2816 171030
## 2    2811 168164  5     2865.4 9.5794 4.453e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
AIC(mod_base, mod_weekend, mod_weekday)
```

```
##          df      AIC
## mod_base    14 19696.29
## mod_weekend 15 19668.61
```

```

## mod_weekday 20 19630.79
# --- Manual Time Series Cross-Validation ---

# 1. Setup
n_total <- length(bike_ts)
min_train <- 365 * 2 # Start training after 2 years of data
horizon <- 1           # One-step ahead forecast

# We will test degrees 1 to 6
degrees <- 1:6
cv_mse <- numeric(length(degrees))

# Time index and Monthly dummies for the WHOLE dataset first
# (It's safe to create predictors beforehand as they are deterministic)
time_idx <- 1:n_total
month_fac <- factor(cycle(bike_ts)) # 1 to 12

# Prepare the master dataframe
data_full <- data.frame(y = as.numeric(bike_ts), t = time_idx, month = month_fac)

# Loop through degrees
for (deg in degrees) {

  errors <- c()

  # 2. Rolling Loop
  # We slide the cutoff "i" from min_train to (n_total - 1)
  # This might be slow if n is large. For 2000 points it's fine.
  # To speed up, you can step by 30 days (monthly rolling) instead of 1 day.

  for (i in seq(min_train, n_total - 1, by = 30)) { # Jumping by 30 days for speed

    # Train set: 1 to i
    train_set <- data_full[1:i, ]

    # Test set: i+1
    test_set <- data_full[(i + 1), , drop = FALSE] # drop=FALSE keeps it a dataframe

    # Fit model on Train
    # Note: raw=TRUE is critical here so the coefficients mean something stable
    fit <- lm(y ~ poly(t, deg, raw = TRUE) + month, data = train_set)

    # Predict on Test
    pred <- predict(fit, newdata = test_set)

    # Calculate Squared Error
    err <- (test_set$y - pred)^2
    errors <- c(errors, err)
  }

  # Average MSE for this degree
  cv_mse[deg] <- mean(errors)
  print(paste("Degree", deg, "MSE:", round(cv_mse[deg], 4)))
}

```

```

}

## [1] "Degree 1 MSE: 15887410.39"
## [1] "Degree 2 MSE: 14393271.5257"
## [1] "Degree 3 MSE: 15023753.3796"
## [1] "Degree 4 MSE: 15601016.1575"
## [1] "Degree 5 MSE: 16379940.025"
## [1] "Degree 6 MSE: 15553902.3576"

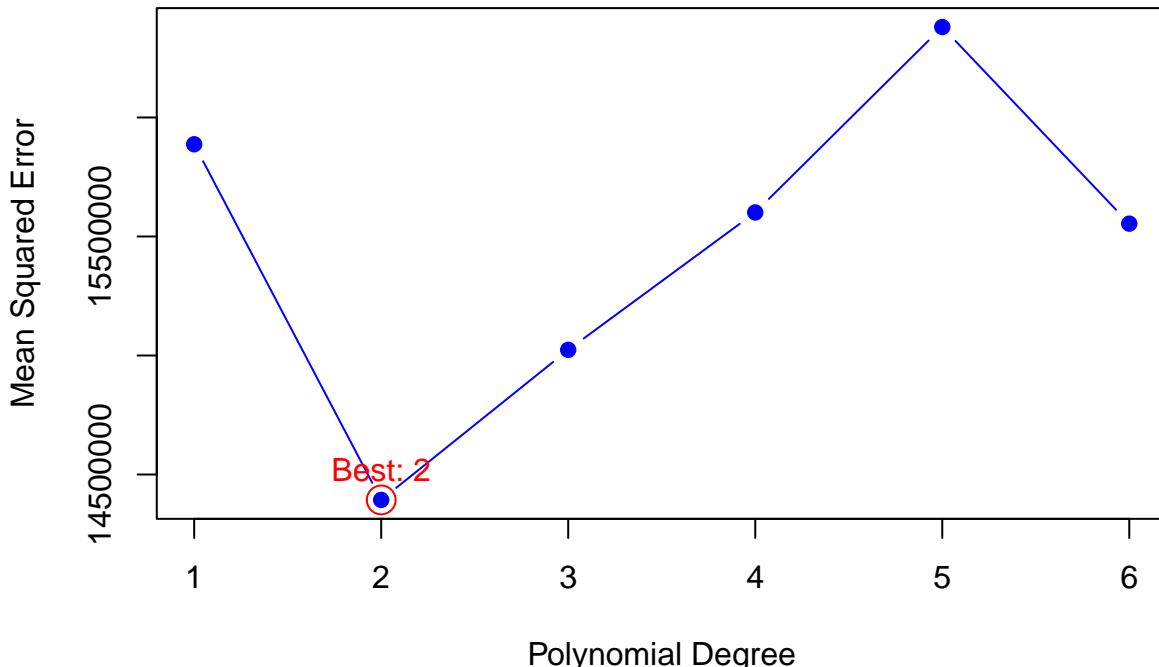
# 3. Visualize Results
results_ts_cv <- data.frame(Degree = degrees, MSE = cv_mse)

plot(results_ts_cv$Degree, results_ts_cv$MSE, type = "b", col = "blue", pch = 19,
      main = "Time Series Cross-Validation Error",
      xlab = "Polynomial Degree", ylab = "Mean Squared Error")

best_d <- which.min(cv_mse)
points(best_d, cv_mse[best_d], col = "red", cex = 2, pch = 1)
text(best_d, cv_mse[best_d], paste("Best:", best_d), pos = 3, col = "red")

```

Time Series Cross–Validation Error



```

# --- Define the range of degrees to test ---
max_degree <- 10 # You can go higher if you want, but 6-8 is usually enough
output <- data.frame(Degree = integer(), AdjR2 = double(), AIC = double(), BIC = double(), CV_MSE = dou

# --- Setup for Cross-Validation ---
set.seed(443) # For reproducibility
n <- nrow(model_data)
k_folds <- 5
# Create random fold assignments
folds <- sample(rep(1:k_folds, length.out = n))

```

```

for (p in 1:max_degree) {

  # 1. Fit model on FULL data to get AIC/BIC/AdjR2
  full_mod <- lm(y_trans ~ poly(tim, p, raw = FALSE) + month + weekday, data = model_data)

  # 2. Cross-Validation Loop
  cv_errors <- numeric(k_folds)
  for (k in 1:k_folds) {
    # Split data
    test_idx <- which(folds == k)
    train_data <- model_data[-test_idx, ]
    test_data <- model_data[test_idx, ]

    # Fit on Train
    train_mod <- lm(y_trans ~ poly(tim, p, raw = FALSE) + month + weekday, data = train_data)

    # Predict on Test
    pred_test <- predict(train_mod, newdata = test_data)

    # Calc MSE
    cv_errors[k] <- mean((test_data$y_trans - pred_test)^2)
  }

  # 3. Store Results
  output[p, ] <- c(p,
                     summary(full_mod)$adj.r.squared,
                     AIC(full_mod),
                     BIC(full_mod),
                     mean(cv_errors))
}

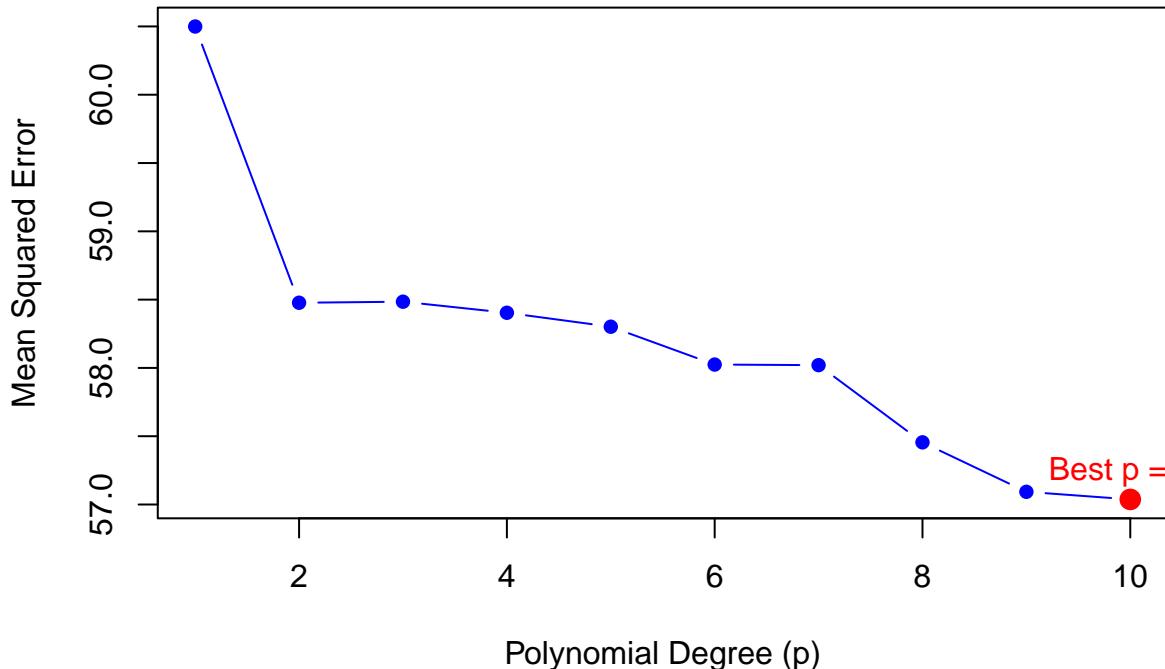
# --- View the Selection Table ---
print(round(output, 4))

##      Degree  AdjR2      AIC      BIC  CV_MSE
## 1      1 0.8394 19630.79 19749.75 60.4994
## 2      2 0.8448 19535.02 19659.93 58.4771
## 3      3 0.8448 19536.66 19667.51 58.4847
## 4      4 0.8449 19534.22 19671.02 58.4037
## 5      5 0.8454 19527.25 19670.00 58.3015
## 6      6 0.8461 19514.63 19663.33 58.0245
## 7      7 0.8462 19514.96 19669.61 58.0203
## 8      8 0.8476 19488.67 19649.27 57.4550
## 9      9 0.8487 19470.42 19636.96 57.0930
## 10     10 0.8488 19469.74 19642.23 57.0373

# --- Plot CV Error to find the minimum ---
plot(output$Degree, output$CV_MSE, type = 'b', pch = 16, col = "blue",
      main = "Cross-Validation Error by Trend Degree",
      xlab = "Polynomial Degree (p)", ylab = "Mean Squared Error")
# Mark the best one
best_p <- which.min(output$CV_MSE)
points(best_p, output$CV_MSE[best_p], col = "red", pch = 16, cex = 1.5)
text(best_p, output$CV_MSE[best_p], labels = paste("Best p =", best_p), pos = 3, col = "red")

```

Cross-Validation Error by Trend Degree



```
# --- APSE with Temporal Split (The "Honest" Test) ---

# 1. Split Data into Training (Past) and Test (Recent Future)
# Let's hold out the last 365 days as our "Test Set"
n_total <- nrow(model_data)
n_test <- 365
n_train <- n_total - n_test

train_data <- model_data[1:n_train, ]
test_data <- model_data[(n_train + 1):n_total, ]

# 2. Loop through degrees and check forecasting error
max_degree <- 8 # Feel free to check up to 8 or 10
apse_results <- numeric(max_degree)

par(mfrow = c(1,1))
# Plot the test data to set up the canvas
plot(test_data$tim, test_data$y_trans, type = "l", col = "black", lwd = 2,
      main = "Forecasting Performance on Hold-out Set (Last 365 Days)",
      xlab = "Time", ylab = "Transformed Trips",
      ylim = range(model_data$y_trans))

colors <- rainbow(max_degree)

for (p in 1:max_degree) {
  # Fit ONLY on Training Data
  mod <- lm(y_trans ~ poly(tim, p, raw = FALSE) + month + weekday, data = train_data)
```

```

# Predict on Test Data
pred <- predict(mod, newdata = test_data)

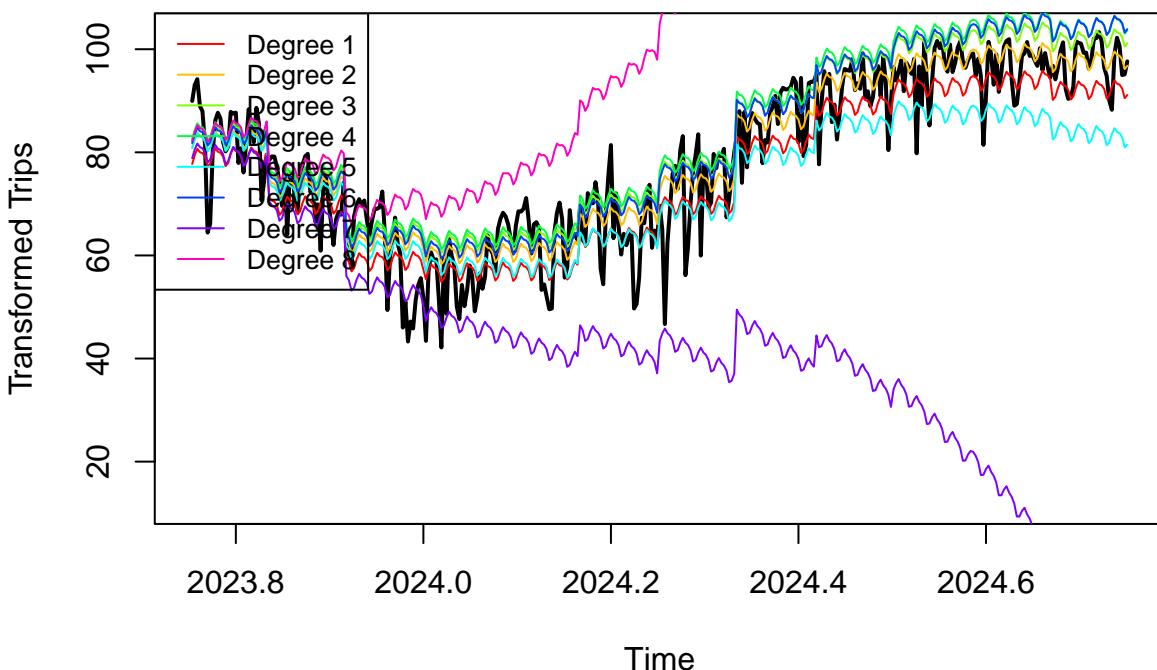
# Calculate APSE (Mean Squared Error on the Test Set)
apse_results[p] <- mean((test_data$y_trans - pred)^2)

# Add line to plot to visualize the "Explosion"
lines(test_data$tim, pred, col = colors[p], lwd = 1)
}

legend("topleft", legend = paste("Degree", 1:max_degree), col = colors, lty = 1, cex = 0.8)

```

Forecasting Performance on Hold-out Set (Last 365 Days)



```

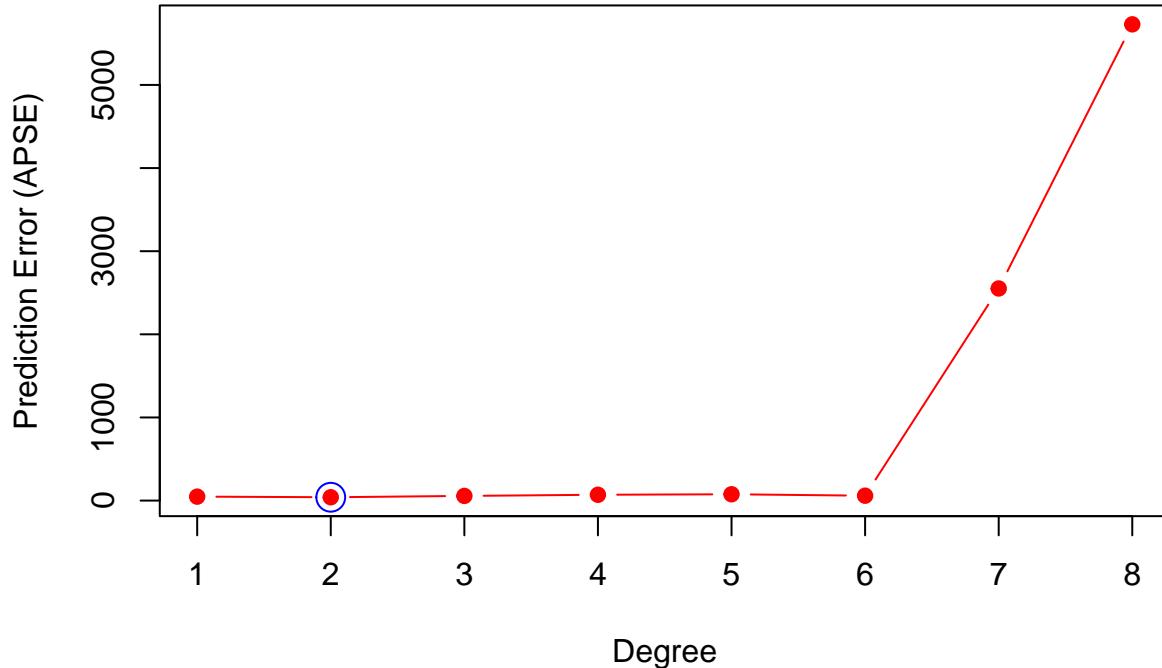
# 3. Print the Results
results_table <- data.frame(Degree = 1:max_degree, APSE = apse_results)
print(results_table)

##   Degree      APSE
## 1      1 47.62278
## 2      2 40.26776
## 3      3 56.12507
## 4      4 69.93335
## 5      5 76.13875
## 6      6 58.66346
## 7      7 2551.62090
## 8      8 5728.81653

# Plot APSE scores
plot(results_table$Degree, results_table$APSE, type = "b", pch = 19, col = "red",
     main = "APSE vs Polynomial Degree", xlab = "Degree", ylab = "Prediction Error (APSE)")
points(which.min(apse_results), min(apse_results), cex = 2, col = "blue")

```

APSE vs Polynomial Degree



```
# --- 1. Refit the Winner (Degree 2) on Full Data ---
final_model <- lm(y_trans ~ poly(tim, 2, raw = FALSE) + month + weekday, data = model_data)
summary(final_model)
```

```
##
## Call:
## lm(formula = y_trans ~ poly(tim, 2, raw = FALSE) + month + weekday,
##      data = model_data)
##
## Residuals:
##    Min     1Q   Median     3Q    Max 
## -37.542 -3.963   0.672   4.740  29.852 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 40.7487   0.5953   68.455 < 2e-16 ***
## poly(tim, 2, raw = FALSE)1 583.3739   7.6311   76.447 < 2e-16 ***
## poly(tim, 2, raw = FALSE)2  76.1305   7.6601   9.939 < 2e-16 ***
## month02       0.6597   0.6993   0.943  0.345588  
## month03       6.3322   0.6829   9.272 < 2e-16 ***
## month04      12.2275   0.6887  17.755 < 2e-16 ***
## month05      23.9512   0.6832  35.060 < 2e-16 ***
## month06      30.6940   0.6890  44.546 < 2e-16 ***
## month07      33.9593   0.6836  49.679 < 2e-16 ***
## month08      34.7098   0.6839  50.755 < 2e-16 ***
## month09      32.8345   0.6898  47.598 < 2e-16 ***
## month10      24.4784   0.7083  34.559 < 2e-16 ***
## month11      14.5496   0.7147  20.357 < 2e-16 ***
## month12       3.1326   0.7087   4.420  1.02e-05 ***
## weekday2      1.8325   0.5347   3.427  0.000618 ***
```

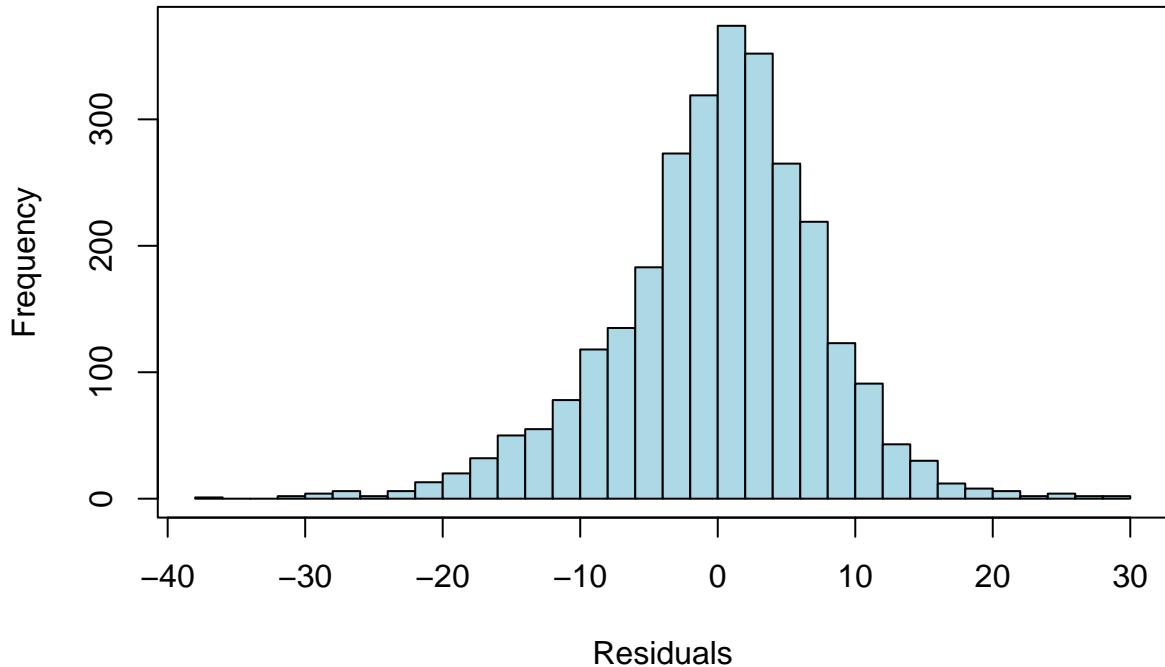
```

## weekday3           2.9754    0.5347    5.565 2.87e-08 ***
## weekday4          2.5131    0.5347    4.700 2.73e-06 ***
## weekday5          2.0453    0.5347    3.825 0.000134 ***
## weekday6          1.0950    0.5347    2.048 0.040666 *
## weekday7         -0.8880    0.5343   -1.662 0.096634 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.603 on 2810 degrees of freedom
## Multiple R-squared:  0.8458, Adjusted R-squared:  0.8448
## F-statistic: 811.4 on 19 and 2810 DF,  p-value: < 2.2e-16
# --- 2. Diagnostic Plots (Generated One by One) ---

# Plot A: Histogram of Residuals
# Checks if the noise is roughly Bell-shaped
hist(final_model$residuals,
      breaks = 30,
      col = "lightblue",
      main = "Histogram of Residuals",
      xlab = "Residuals")
box() # Adds a nice border

```

Histogram of Residuals

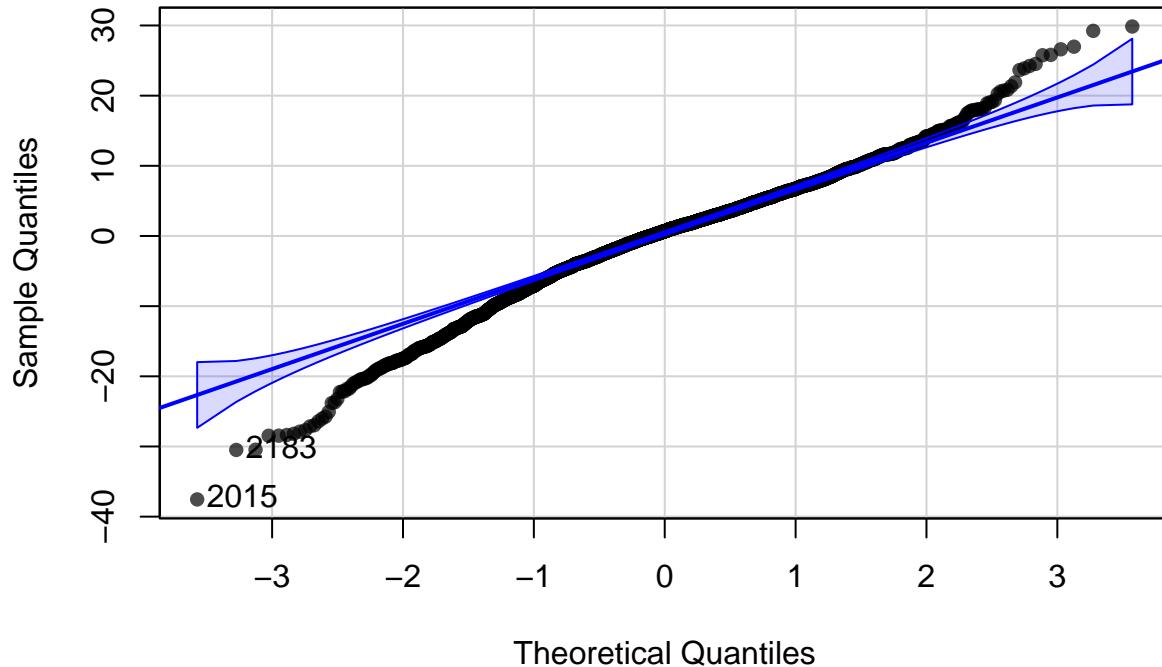


```

# Plot B: Q-Q Plot
# Checks for Normality (points should hug the blue line)
car::qqPlot(final_model$residuals,
            pch = 16,
            col = adjustcolor("black", 0.7),
            main = "Normal Q-Q Plot",
            xlab = "Theoretical Quantiles",
            ylab = "Sample Quantiles")

```

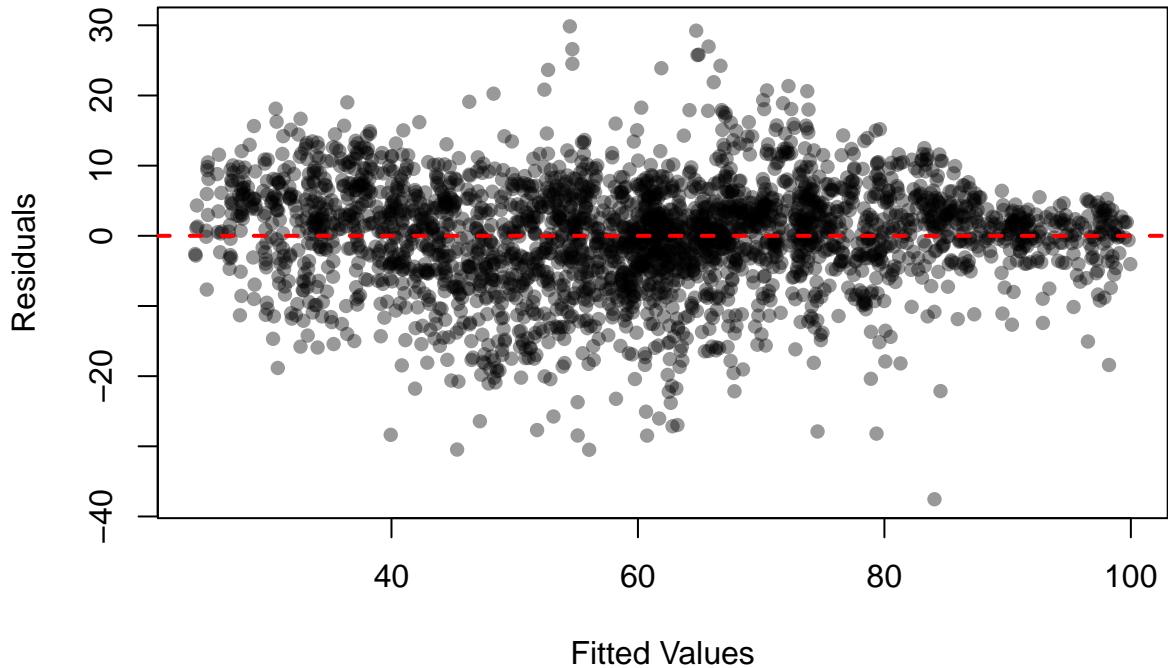
Normal Q–Q Plot



```
## [1] 2015 2183
```

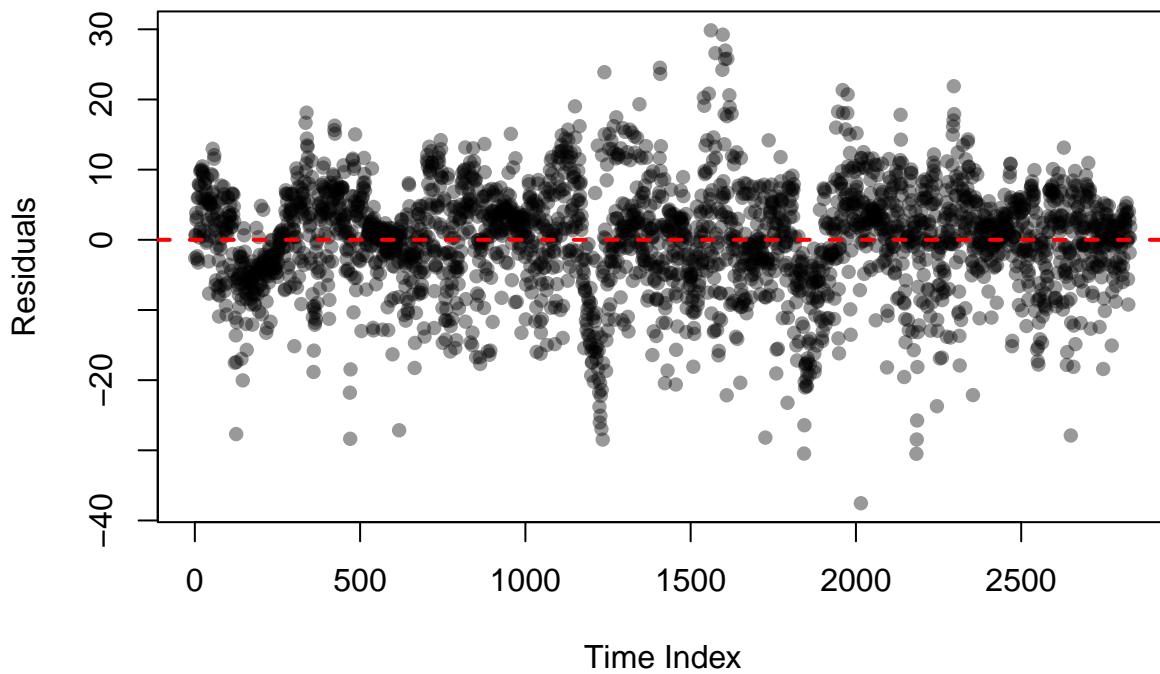
```
# Plot C: Fitted Values vs. Residuals
# Checks for Homoscedasticity (spread should be constant, no funnel shape)
plot(fitted(final_model), residuals(final_model),
      pch = 16,
      col = adjustcolor("black", 0.4),
      main = "Residuals vs. Fitted Values",
      xlab = "Fitted Values",
      ylab = "Residuals")
abline(h = 0, lty = 2, col = "red", lwd = 2)
```

Residuals vs. Fitted Values



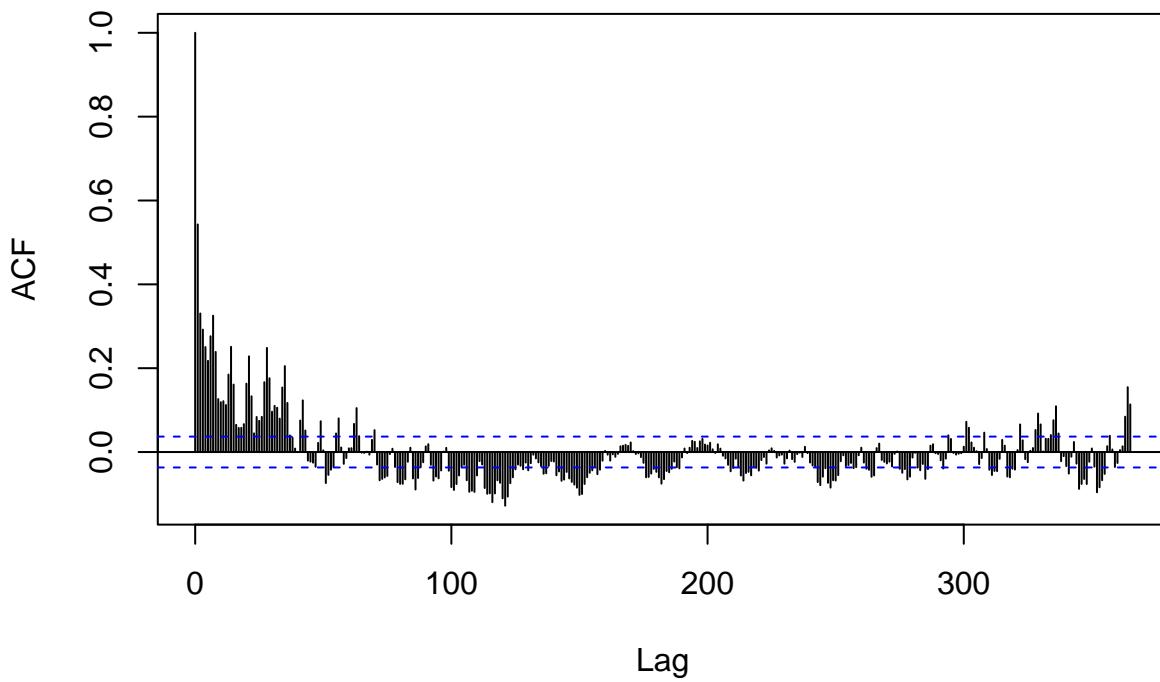
```
# Plot D: Residuals vs. Time
# Checks if variance changes over time or if there are bursts of volatility
plot(residuals(final_model),
      pch = 16,
      col = adjustcolor("black", 0.4),
      main = "Residuals vs. Time",
      xlab = "Time Index",
      ylab = "Residuals")
abline(h = 0, lty = 2, col = "red", lwd = 2)
```

Residuals vs. Time



```
# Plot E: ACF of Residuals
# Checks for Independence (bars should stay inside blue lines)
acf(residuals(final_model),
  lag.max = 365, # Check up to a year of lags to see long-term correlation
  main = "ACF of Residuals")
```

ACF of Residuals



```

# --- 3. Statistical Tests ---
print("--- Shapiro-Wilk Test for Normality ---")

## [1] "--- Shapiro-Wilk Test for Normality ---"
# Limit to 5000 samples because shapiro.test() has a hard limit
if(length(residuals(final_model)) > 5000) {
  print(shapiro.test(residuals(final_model)[1:5000]))
} else {
  print(shapiro.test(residuals(final_model)))
}

## 
## Shapiro-Wilk normality test
##
## data: residuals(final_model)
## W = 0.98126, p-value < 2.2e-16
print("--- Kolmogorov-Smirnov Test for Normality ---")

## [1] "--- Kolmogorov-Smirnov Test for Normality ---"
print(ks.test(scale(residuals(final_model)), "pnorm"))

## 
## Asymptotic one-sample Kolmogorov-Smirnov test
##
## data: scale(residuals(final_model))
## D = 0.052341, p-value = 3.689e-07
## alternative hypothesis: two-sided
print("--- Runs Test for Randomness ---")

## [1] "--- Runs Test for Randomness ---"
# Checks if residuals have random sign flips
print(randtests::runs.test(residuals(final_model)))

## 
## Runs Test
##
## data: residuals(final_model)
## statistic = -22.561, runs = 816, n1 = 1415, n2 = 1415, n = 2830,
## p-value < 2.2e-16
## alternative hypothesis: nonrandomness

```