

Step 1: Loading in the Data

```
# Packages used in tutorials
library(MASS)      # boxcox
library(car)        # qqPlot

## Loading required package: carData
library(randtests) # runs.test

## Warning: package 'randtests' was built under R version 4.3.3
# library(forecast) # OPTIONAL if you want auto.arima, not required

bike <- read.csv("trips_per_day.csv")
bike$trip_date <- as.Date(bike$trip_date)

str(bike)

## 'data.frame': 2969 obs. of 2 variables:
##   $ trip_date: Date, format: "2016-01-10" "2016-01-11" ...
##   $ n_trips  : int 2273 3623 2535 2966 2970 2636 4122 3104 1642 4834 ...
head(bike)

##    trip_date n_trips
## 1 2016-01-10     2273
## 2 2016-01-11     3623
## 3 2016-01-12     2535
## 4 2016-02-10     2966
## 5 2016-02-11     2970
## 6 2016-02-12     2636

range(bike$trip_date)

## [1] "2016-01-10" "2024-09-30"

Initial Plotting for Time Series

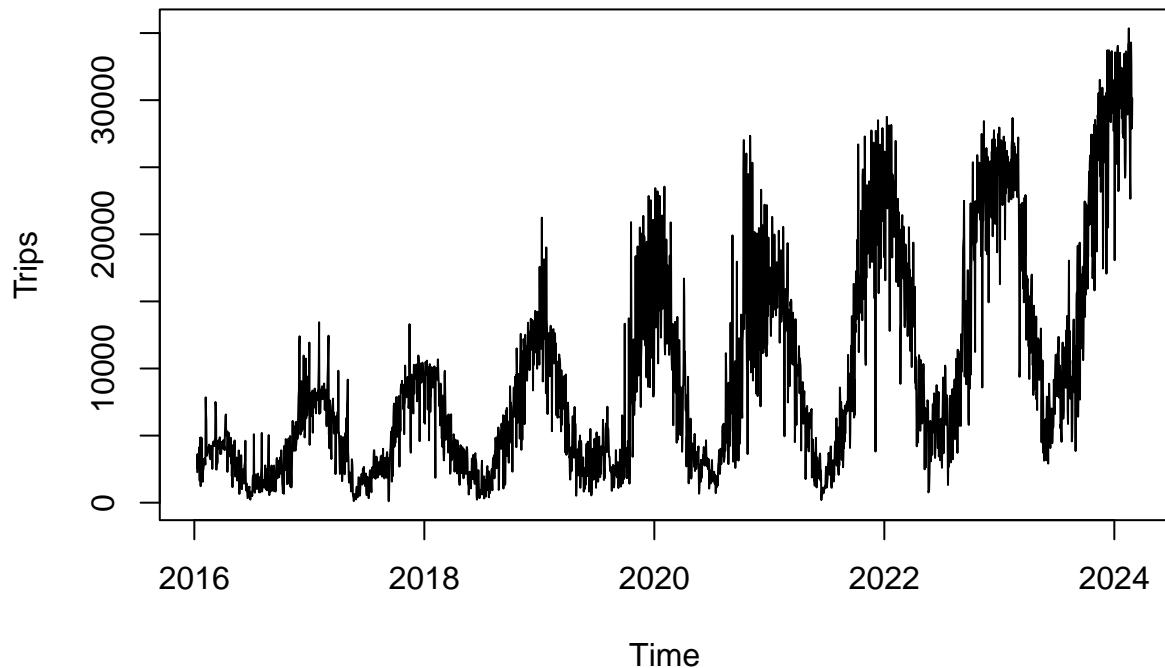
# Sort just in case
bike <- bike[order(bike$trip_date), ]

# Extract response as a vector
y <- bike$n_trips

# Daily frequency with yearly seasonality (approx 365)
bike_ts <- ts(
  y,
  start = c(as.numeric(format(min(bike$trip_date), "%Y")),
            as.numeric(format(min(bike$trip_date), "%j"))),
  frequency = 365
)

plot(bike_ts, main = "Daily BikeShare Trips in Toronto", ylab = "Trips")
```

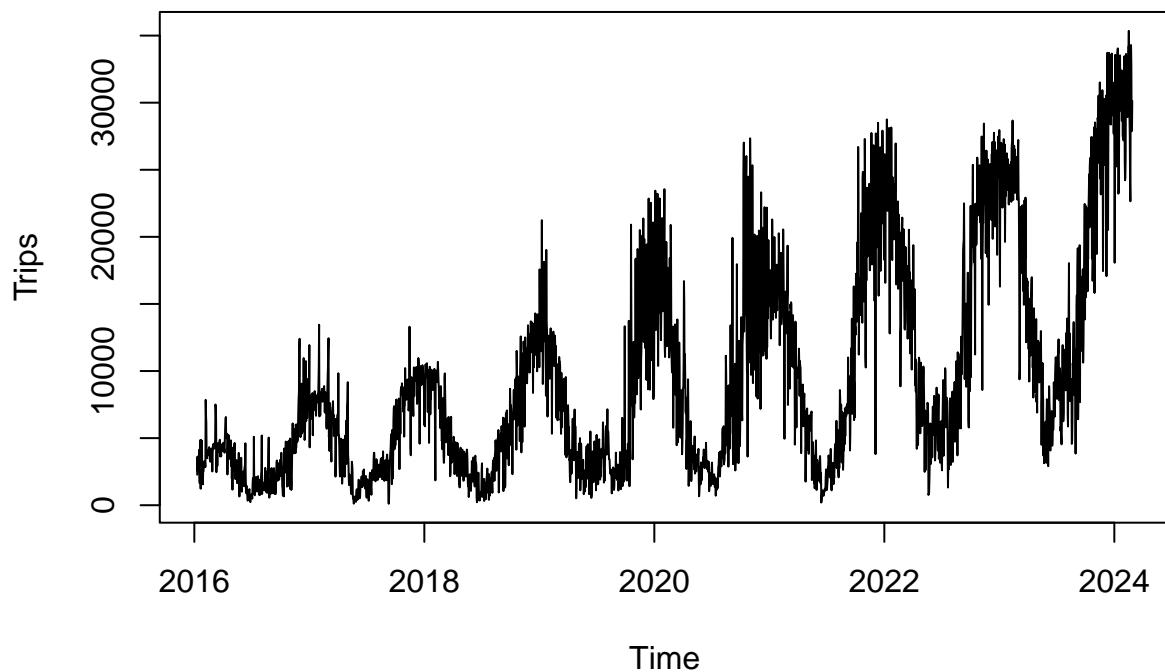
Daily BikeShare Trips in Toronto



Step 2: EDA

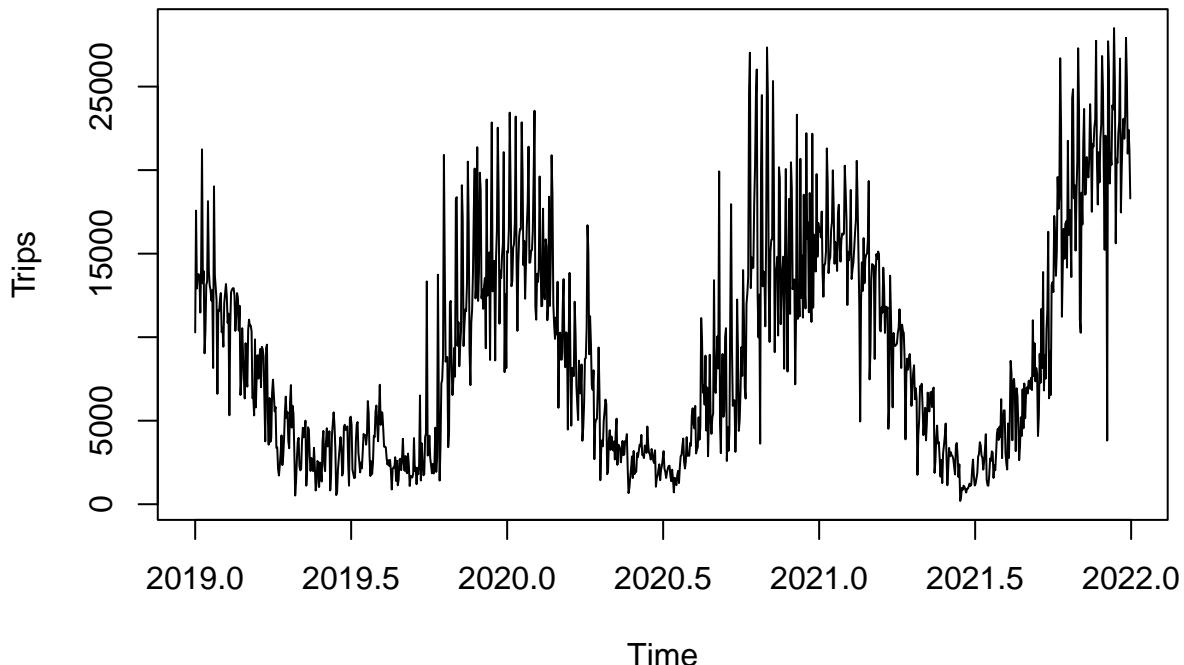
```
par(mfrow = c(1, 1))
plot(bike_ts, main = "Daily Trips", ylab = "Trips")
```

Daily Trips



```
# maybe a zoom on a couple of years
plot(window(bike_ts, start = c(2019, 1), end = c(2021, 365)),
     main = "Daily Trips: 2019-2021", ylab = "Trips")
```

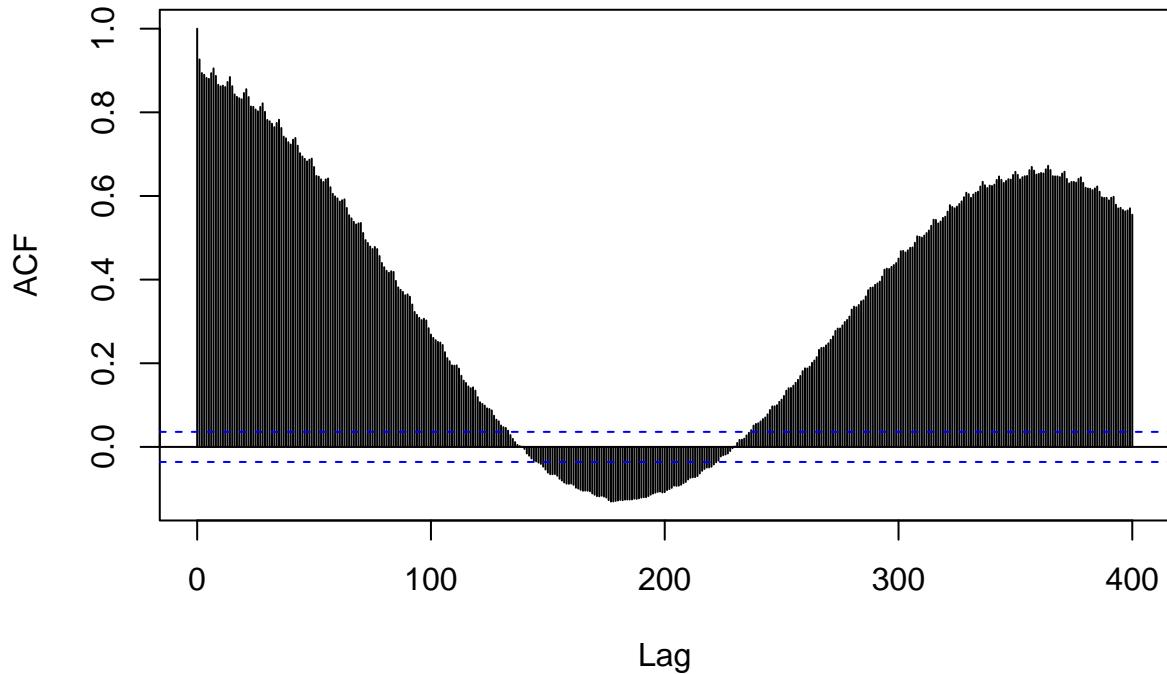
Daily Trips: 2019-2021



Confirming Seasonality with ACF and Spectrum

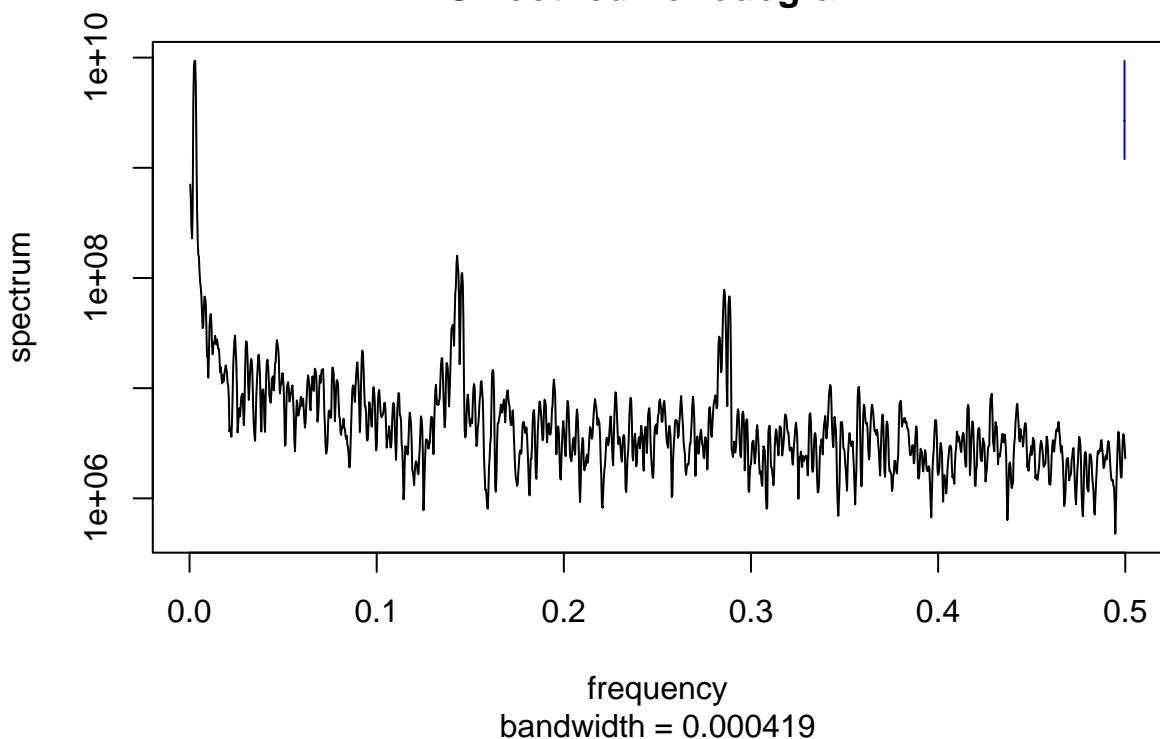
```
acf(as.vector(bike_ts), lag.max = 400,
    main = "ACF of Daily Trips")
```

ACF of Daily Trips



```
spec_bike <- spectrum(as.vector(bike_ts), spans = 5)
```

Series: x
Smoothed Periodogram

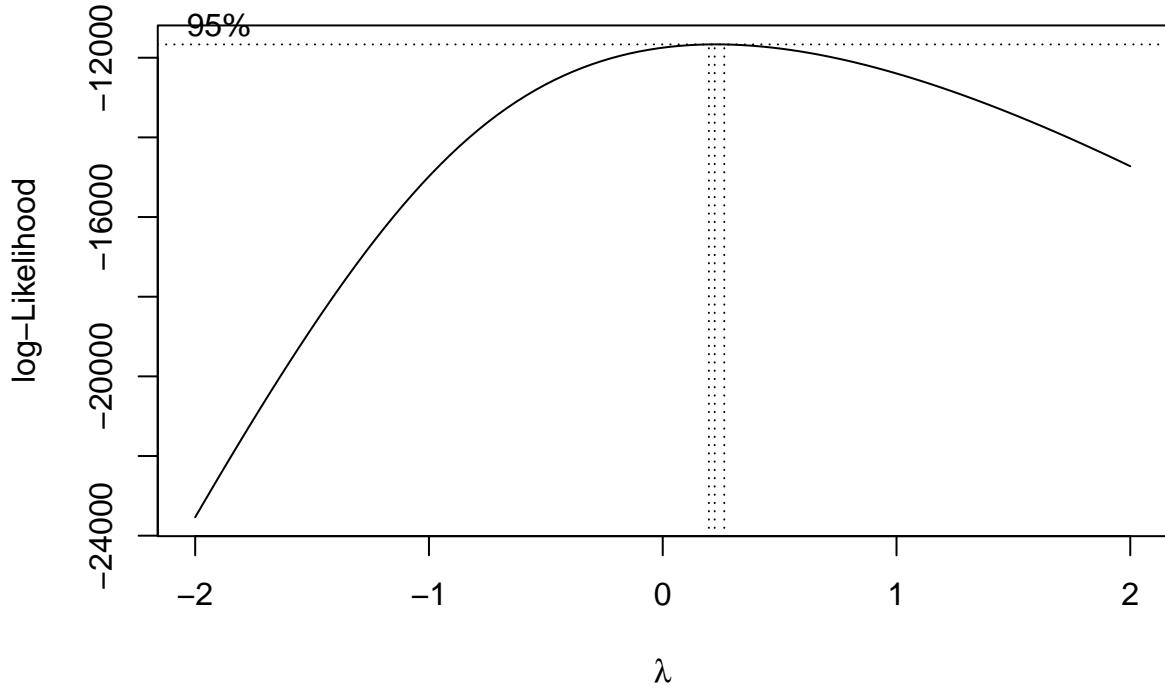


```
1 / spec_bike$freq[which.max(spec_bike$spec)] # estimated period
```

```
## [1] 333.3333
```

Step 3: Box-Cox transformation

```
# Simple intercept-only model (like in tutorial)
bc_model_raw <- lm(bike_ts ~ 1)
boxcox_raw <- MASS::boxcox(bc_model_raw, lambda = seq(-2, 2, 0.1))
```



```
(lambda_opt_raw <- boxcox_raw$x[which.max(boxcox_raw$y)])
```

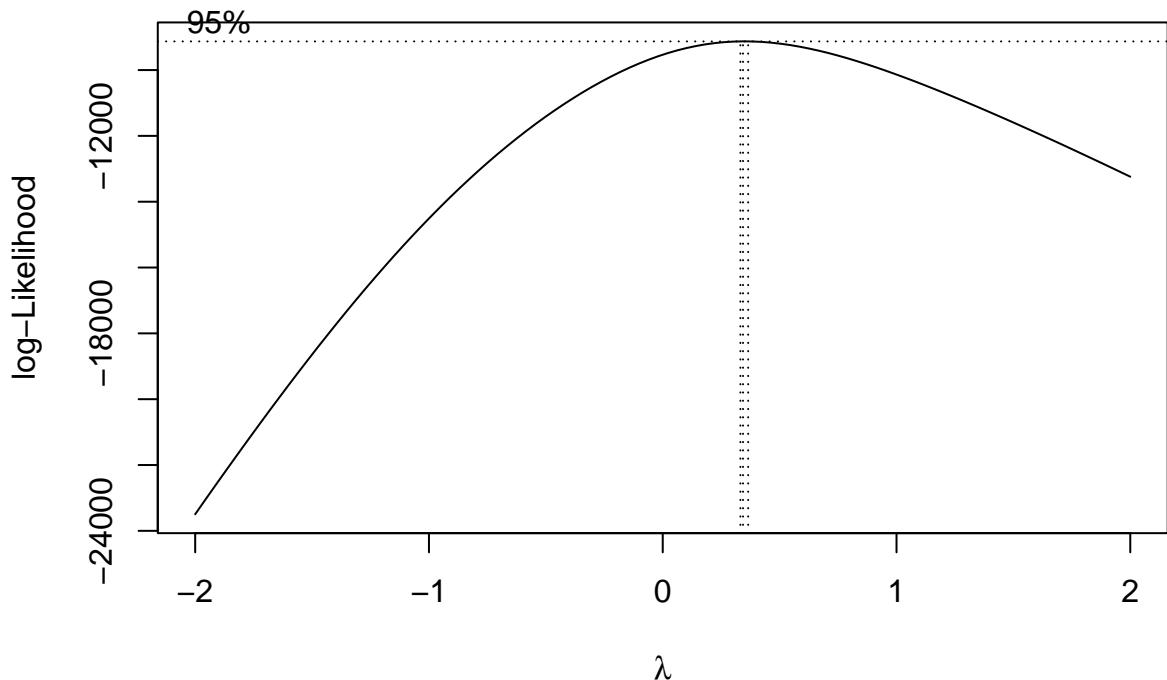
```
## [1] 0.2222222
```

```
tim <- time(bike_ts) # continuous time index
```

```
# Season: year and day-of-year or month; simplest is month
# Build a monthly factor from dates (instead of cycle, since this is daily)
month <- factor(format(bike$trip_date, "%m"))
```

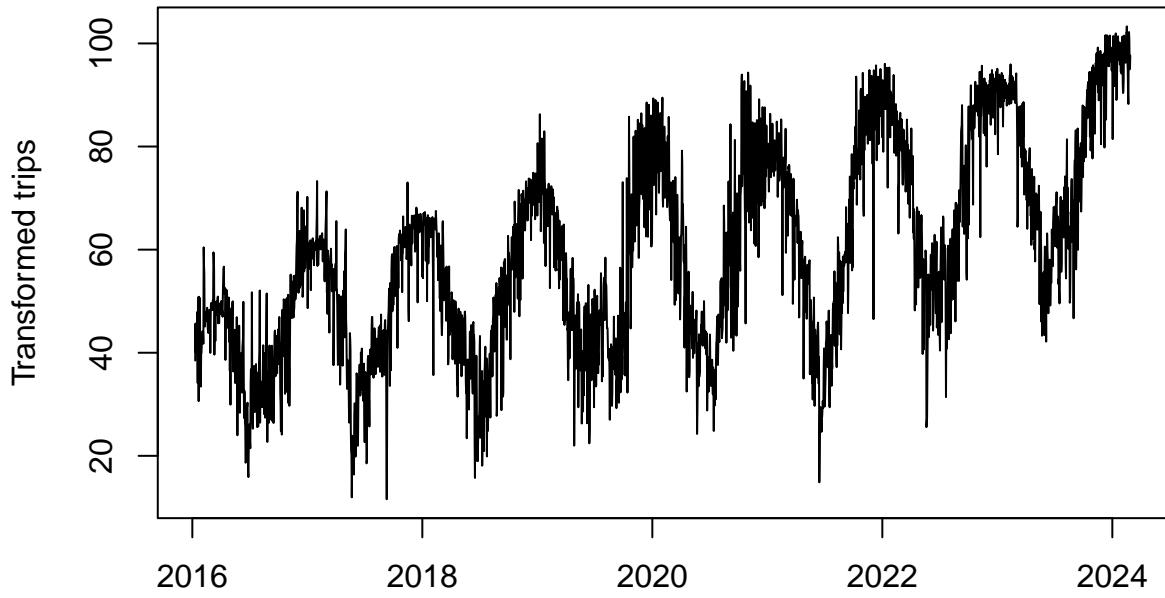
```
reg_for_bc <- lm(bike_ts ~ tim + month)
```

```
boxcox_mod <- MASS::boxcox(reg_for_bc, lambda = seq(-2, 2, 0.1))
```



```
(lambda_opt_mod <- boxcox_mod$x[which.max(boxcox_mod$y)])  
  
## [1] 0.3434343  
lam <- lambda_opt_mod # keep this for later  
  
if (lam == 0) {  
  y_trans <- log(bike_ts)  
} else if (lam > 0) {  
  y_trans <- (bike_ts^lam - 1) / lam  
} else {  
  # negative lambda + use minus sign trick like in lectures  
  y_trans <- -(bike_ts^lam)  
}  
  
plot(y_trans, main = "Transformed Daily Trips", ylab = "Transformed trips")
```

Transformed Daily Trips

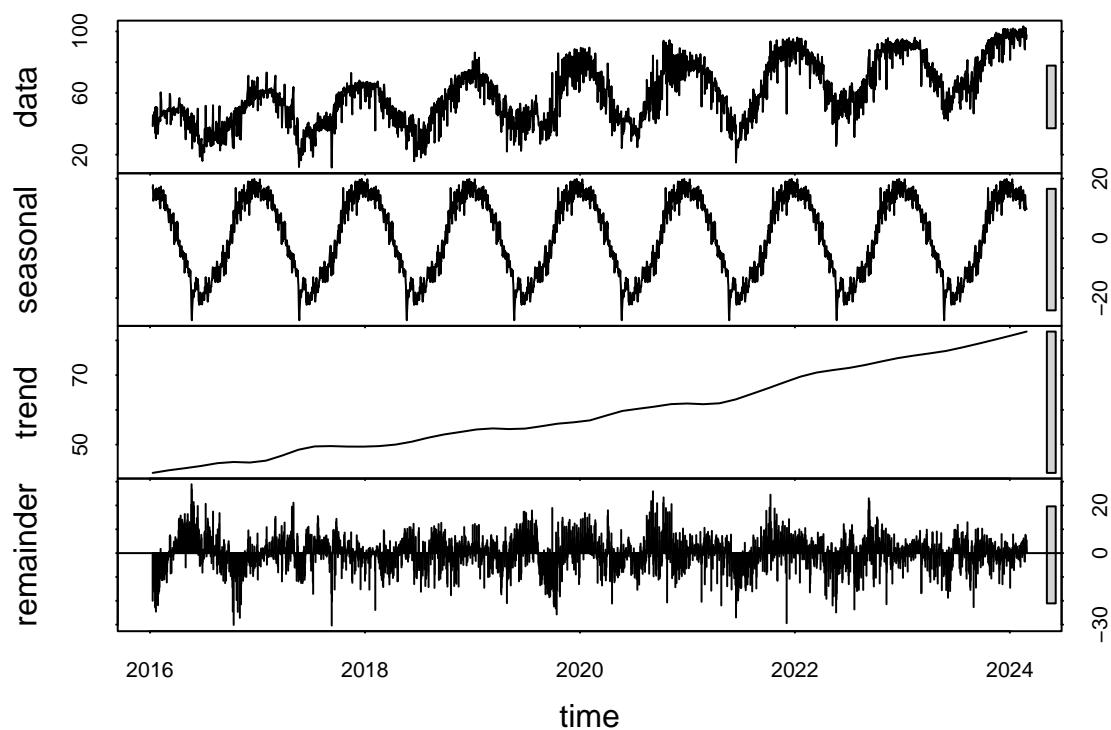


cal Decomposition

```
Decomp_bike <- stl(y_trans, s.window = "periodic")
plot(Decomp_bike, main = "STL Decomposition of Transformed Daily Trips")
```

Classi-

STL Decomposition of Transformed Daily Trips



```

# Extract components if needed
bike_seasonal <- Decomps_bike$time.series[, "seasonal"]
bike_trend     <- Decomps_bike$time.series[, "trend"]
bike_remainder <- Decomps_bike$time.series[, "remainder"]

tim   <- time(y_trans)
month <- factor(format(bike$trip_date, "%m"))
dow   <- factor(weekdays(bike$trip_date)) # optional, but nice

# Simple model: linear trend + monthly seasonality
reg1 <- lm(y_trans ~ tim + month)
summary(reg1)

##
## Call:
## lm(formula = y_trans ~ tim + month)
##
## Residuals:
##    Min      1Q  Median      3Q      Max
## -36.930 -4.421   0.841   5.135  33.503
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.010e+04  1.288e+02 -78.419 < 2e-16 ***
## tim          5.022e+00  6.378e-02  78.738 < 2e-16 ***
## month02      6.976e-01  7.537e-01   0.926  0.354735
## month03      6.327e+00  7.358e-01   8.599 < 2e-16 ***
## month04      1.183e+01  7.421e-01  15.935 < 2e-16 ***
## month05      2.364e+01  7.360e-01  32.113 < 2e-16 ***
## month06      3.001e+01  7.424e-01  40.431 < 2e-16 ***
## month07      3.270e+01  7.130e-01  45.862 < 2e-16 ***
## month08      3.376e+01  7.130e-01  47.344 < 2e-16 ***
## month09      3.207e+01  7.186e-01  44.635 < 2e-16 ***
## month10      2.345e+01  7.404e-01  31.666 < 2e-16 ***
## month11      1.429e+01  7.468e-01  19.132 < 2e-16 ***
## month12      2.715e+00  7.403e-01   3.667  0.000249 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.143 on 2956 degrees of freedom
## Multiple R-squared:  0.8212, Adjusted R-squared:  0.8205
## F-statistic: 1131 on 12 and 2956 DF, p-value: < 2.2e-16

reg2 <- lm(y_trans ~ poly(tim, 2, raw = FALSE) + month)
summary(reg2)

##
## Call:
## lm(formula = y_trans ~ poly(tim, 2, raw = FALSE) + month)
##
## Residuals:
##    Min      1Q  Median      3Q      Max
## -36.363 -4.468   0.913   5.123  33.023
##
## Coefficients:
```

```

##                               Estimate Std. Error t value Pr(>|t|)
## (Intercept)                 41.6116   0.5175 80.412 < 2e-16 ***
## poly(tim, 2, raw = FALSE)1 642.5618   8.1150 79.182 < 2e-16 ***
## poly(tim, 2, raw = FALSE)2 47.9377   8.1680  5.869 4.87e-09 ***
## month02                      0.6949   0.7495  0.927 0.353935
## month03                      6.3211   0.7317  8.639 < 2e-16 ***
## month04                     11.8129   0.7379 16.008 < 2e-16 ***
## month05                     23.6142   0.7319 32.264 < 2e-16 ***
## month06                     29.9821   0.7382 40.614 < 2e-16 ***
## month07                     32.4508   0.7102 45.690 < 2e-16 ***
## month08                     33.5098   0.7102 47.183 < 2e-16 ***
## month09                     31.8278   0.7158 44.468 < 2e-16 ***
## month10                     23.4850   0.7363 31.897 < 2e-16 ***
## month11                     14.3325   0.7427 19.299 < 2e-16 ***
## month12                     2.7584   0.7362  3.747 0.000183 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.097 on 2955 degrees of freedom
## Multiple R-squared:  0.8233, Adjusted R-squared:  0.8225
## F-statistic:  1059 on 13 and 2955 DF,  p-value: < 2.2e-16
reg3 <- lm(y_trans ~ poly(tim, 2, raw = FALSE) + month + dow)
summary(reg3)

```

```

##
## Call:
## lm(formula = y_trans ~ poly(tim, 2, raw = FALSE) + month + dow)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -37.086  -4.311   0.836   4.901  31.512 
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)
## (Intercept)                 42.3325   0.6270 67.516 < 2e-16 ***
## poly(tim, 2, raw = FALSE)1 642.5699   8.0310 80.011 < 2e-16 ***
## poly(tim, 2, raw = FALSE)2 48.1134   8.0835  5.952 2.96e-09 ***
## month02                      0.6592   0.7418  0.889 0.374229
## month03                      6.2757   0.7242  8.666 < 2e-16 ***
## month04                     11.8192   0.7303 16.184 < 2e-16 ***
## month05                     23.5691   0.7244 32.538 < 2e-16 ***
## month06                     29.9529   0.7307 40.995 < 2e-16 ***
## month07                     32.4546   0.7029 46.171 < 2e-16 ***
## month08                     33.4604   0.7029 47.603 < 2e-16 ***
## month09                     31.8053   0.7084 44.897 < 2e-16 ***
## month10                     23.4927   0.7287 32.240 < 2e-16 ***
## month11                     14.2872   0.7350 19.438 < 2e-16 ***
## month12                      2.7472   0.7287  3.770 0.000166 ***
## dowMonday                   -1.9523   0.5495 -3.553 0.000387 ***
## dowSaturday                  -1.0741   0.5501 -1.953 0.050951 .
## dowSunday                    -2.8167   0.5501 -5.120 3.25e-07 ***
## dowThursday                  0.3757   0.5501  0.683 0.494652
## dowTuesday                  -0.2132   0.5501 -0.388 0.698380
## dowWednesday                 0.8004   0.5508  1.453 0.146250

```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.014 on 2949 degrees of freedom
## Multiple R-squared:  0.8273, Adjusted R-squared:  0.8262
## F-statistic: 743.3 on 19 and 2949 DF,  p-value: < 2.2e-16
```